

```
In [34]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
import seaborn as sns
```

```
In [35]: df = pd.read_csv("data.csv")
df.columns.to_list()
```

```
Out[35]: ['person_age',
'person_income',
'person_home_ownership',
'person_emp_length',
'loan_intent',
'loan_grade',
'loan_amnt',
'loan_int_rate',
'loan_status',
'loan_percent_income',
'cb_person_default_on_file',
'cb_person_cred_hist_length']
```

## Features

- person\_age: Age of the individual applying for the loan.
- person\_income: Annual income of the individual.
- person\_home\_ownership: Type of home ownership of the individual.
- rent: The individual is currently renting a property.
- mortgage: The individual has a mortgage on the property they own.
- own: The individual owns their home outright.
- other: Other categories of home ownership that may be specific to the dataset.
- person\_emp\_length: Employment length of the individual in years.
- loan\_intent: The intent behind the loan application.
- loan\_grade: The grade assigned to the loan based on the creditworthiness of the borrower.
  - A: The borrower has a high creditworthiness, indicating low risk.
  - B: The borrower is relatively low-risk, but not as creditworthy as Grade A.
  - C: The borrower's creditworthiness is moderate.
  - D: The borrower is considered to have higher risk compared to previous grades.
  - E: The borrower's creditworthiness is lower, indicating a higher risk.
  - F: The borrower poses a significant credit risk.
  - G: The borrower's creditworthiness is the lowest, signifying the highest risk.
- loan\_amnt: The loan amount requested by the individual.
- loan\_int\_rate: The interest rate associated with the loan.
- loan\_status: Loan status, where 0 indicates non-default and 1 indicates default.
  - 0: Non-default - The borrower successfully repaid the loan as agreed, and there was no default.

- 1: Default - The borrower failed to repay the loan according to the agreed-upon terms and defaulted on the loan.
- loan\_percent\_income: The percentage of income represented by the loan amount.
- cb\_person\_default\_on\_file: Historical default of the individual as per credit bureau records.
  - Y: The individual has a history of defaults on their credit file.
  - N: The individual does not have any history of defaults.
- cb\_person\_cred\_hist\_length: The length of credit history for the individual.

## Hypothesis

- Given a factors, is it possible to predict if an individual will default on their credit?  
aka target = loan\_status

Extension: what are the features that influence whether someone will default?

```
In [36]: # Do EDA to understand the data
df.describe()
# have to clean person_age, person_emp_length, loan_pct_income = 0?
```

```
Out [36]:
```

	person_age	person_income	person_emp_length	loan_amnt	loan_int
<b>count</b>	32581.000000	3.258100e+04	31686.000000	32581.000000	29465.00
<b>mean</b>	27.734600	6.607485e+04	4.789686	9589.371106	11.00
<b>std</b>	6.348078	6.198312e+04	4.142630	6322.086646	3.24
<b>min</b>	20.000000	4.000000e+03	0.000000	500.000000	5.42
<b>25%</b>	23.000000	3.850000e+04	2.000000	5000.000000	7.90
<b>50%</b>	26.000000	5.500000e+04	4.000000	8000.000000	10.99
<b>75%</b>	30.000000	7.920000e+04	7.000000	12200.000000	13.47
<b>max</b>	144.000000	6.000000e+06	123.000000	35000.000000	23.22

## Check for null values and duplicates

There are null values - we either handle them by deletion or imputation later. As for duplicated rows they can be removed first.

```
In [37]: df.isnull().sum()
```


```
Out[37]: person_age      0
         person_income  0
         person_home_ownership  0
         person_emp_length  895
         loan_intent      0
         loan_grade      0
         loan_amnt      0
         loan_int_rate    3116
         loan_status      0
         loan_percent_income  0
         cb_person_default_on_file  0
         cb_person_cred_hist_length  0
         dtype: int64
```

```
In [38]: print(df.duplicated().sum())
         df = df.drop_duplicates()
```

165

```
In [39]: print(df["loan_status"].value_counts(normalize=True))
         # 21% of the loans are defaulted
         plt.figure(figsize=(6, 4))
         df["loan_status"].value_counts().plot(kind="bar")
         plt.title("Loan Default Distribution")
         plt.xlabel("Default Status (1 = Default)")
         plt.ylabel("Count")
         plt.show()
```

```
loan_status
0    0.781312
1    0.218688
Name: proportion, dtype: float64
```



```
In [40]: """
         'person_age', - Numerical
         'person_income', - Numerical
         'person_home_ownership', - Categorical
```

```

'person_emp_length', - Numerical
'loan_intent', - Categorical
'loan_grade', - Categorical
'loan_amnt', - Numerical
'loan_int_rate', - Numerical
'loan_status', - Target
'loan_percent_income', - Numerical
'cb_person_default_on_file', - Categorical
'cb_person_cred_hist_length' - Numerical
"""















# Numerical features
num_features = [
    "person_age",
    "person_income",
    "person_emp_length",
    "loan_amnt",
    "loan_int_rate",
    "loan_percent_income",
    "cb_person_cred_hist_length",
]


fig, axes = plt.subplots(len(num_features), 2, figsize=(15, 4 * len(num_f
for i, feature in enumerate(num_features):
    # Box plot for overall distribution
    sns.boxplot(y=df[feature], ax=axes[i, 0])
    axes[i, 0].set_title(f"Distribution of {feature}")
    axes[i, 0].set_ylabel(feature)

    # Boxplot by target
    sns.boxplot(x="loan_status", y=feature, data=df, ax=axes[i, 1])
    axes[i, 1].set_title(f"{feature} by Default Status")
    axes[i, 1].set_xlabel("Default Status (0=No, 1=Yes)")
    axes[i, 1].set_ylabel(feature)


plt.tight_layout()
plt.show()

```





```
In [41]: plt.figure(figsize=(12, 10))
corr = df[num_features + ["loan_status"]].corr()
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix")
plt.show()
```



INSIGHTS: We thought person\_age and person\_employment\_length would've had a higher correlation, since age and credit history length has a very strong correlation. We investigate this below.

The higher the income, the lower the loan expressed as % of income. The higher the income, the higher the loan amount might be

The higher the loan amount, the higher they were as % of income, i.e. people with lower income take loans that result in higher leverage (to buy big ticket items)

Interest rates show almost no correlation to loan profile which is true in practice since they are driven by mkt forces

Of special interest to us, loan status (last row) has **highest correlation magnitude with loan\_perent\_income and loan\_int\_rate and person\_income.**


```
In [42]: # Visualize relationship between age and employment length

# Scatter plot with color by loan status, turned into function since it w
def drawGraph(df):
    plt.figure(figsize=(12, 6))
    scatter = plt.scatter(
        df["person_age"],
        df["person_emp_length"],
        c=df["loan_status"],
        alpha=0.5,
        cmap="coolwarm",
    )

    plt.colorbar(scatter, label="Loan Status (1=Default)")
    plt.xlabel("Age")
    plt.ylabel("Employment Length (Years)")
    plt.title("Relationship Between Age and Employment Length")

    plt.grid(True, alpha=0.3)
    plt.tight_layout()
    plt.show()

drawGraph(df)
```



## Anomaly Cleaning:


It seemed weird that the correlation between employment length is so weak, so we visualized it. As expected, the correlation is affected by the presence of anomalous data at the top right. Upon closer examination, having an employment length exceeding age is also anomalous, so we removed that as well. This allows our data to fall within a much more reasonable spread.

```
In [43]: # Remove outliers
df["person_emp_length"] = df["person_emp_length"].clip(upper=45)
df["person_age"] = df["person_age"].clip(upper=85)
plt.figure(figsize=(12, 6))

# Remove outliers
df = df[df["person_emp_length"] < df["person_age"]]

drawGraph(df)
```

<Figure size 1200x600 with 0 Axes>



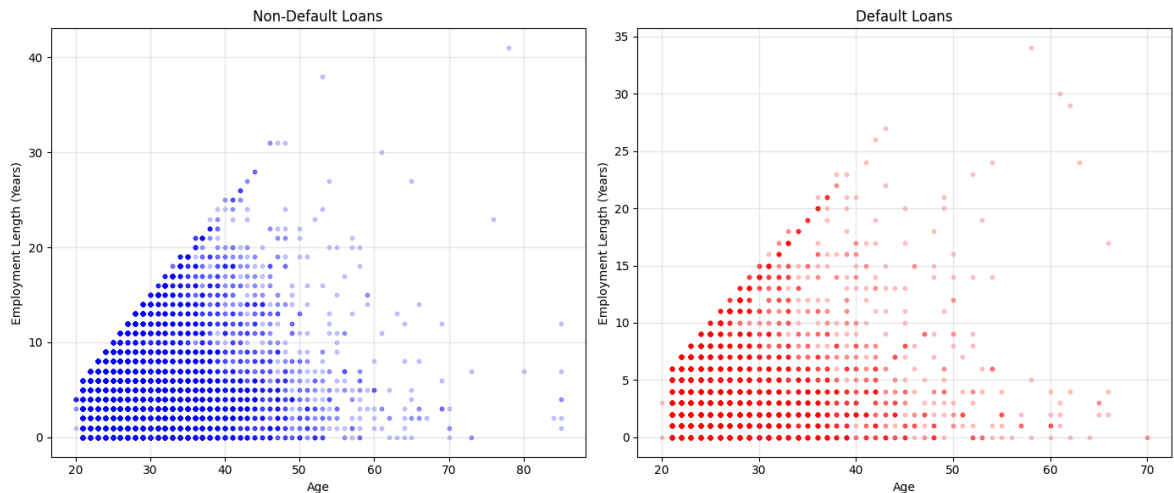
```
In [44]: # Visualize the distribution of the two features in default VS non-default
plt.figure(figsize=(14, 6))

plt.subplot(1, 2, 1)
plt.scatter(
    df[df["loan_status"] == 0]["person_age"],
    df[df["loan_status"] == 0]["person_emp_length"],
    alpha=0.2,
    color="blue",
    s=10,
)
plt.xlabel("Age")
plt.ylabel("Employment Length (Years)")
plt.title("Non-Default Loans")
plt.grid(True, alpha=0.3)

plt.subplot(1, 2, 2)
plt.scatter(
    df[df["loan_status"] == 1]["person_age"],
    df[df["loan_status"] == 1]["person_emp_length"],
    alpha=0.2,
    color="red",
    s=10,
)
plt.xlabel("Age")
plt.ylabel("Employment Length (Years)")
plt.title("Default Loans")
plt.grid(True, alpha=0.3)
```



```
plt.tight_layout()
plt.show()
```

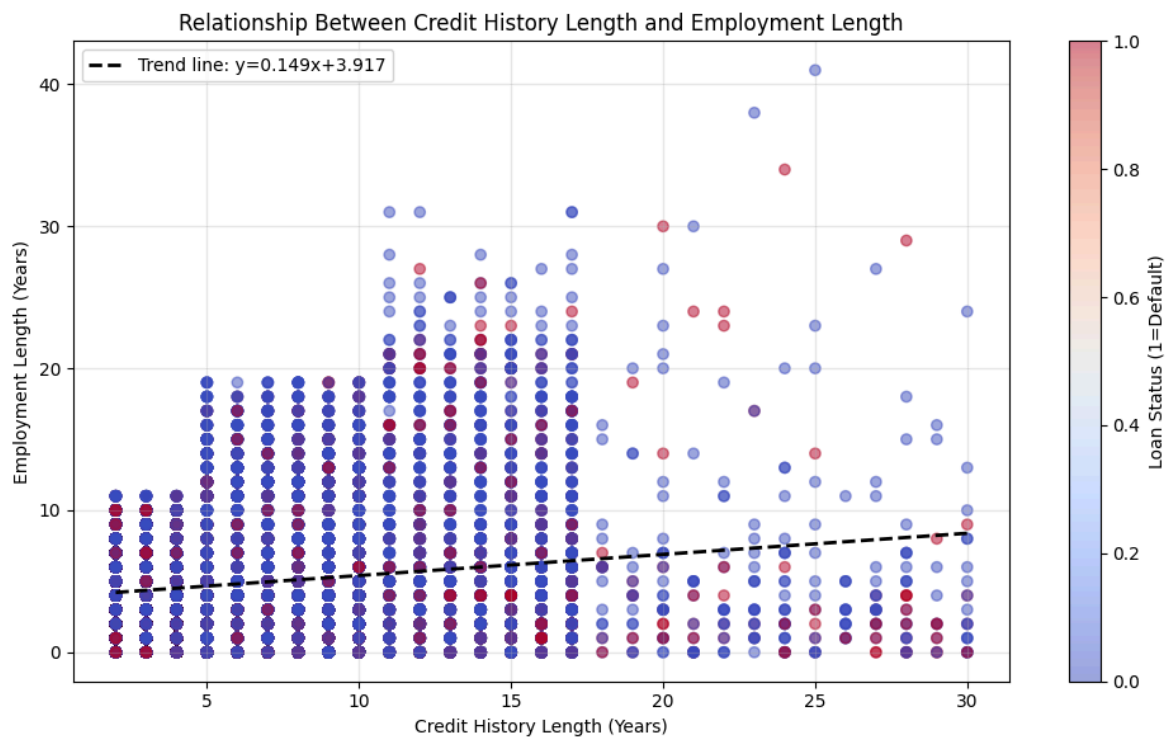


```
In [45]: # Visualize the relationship between credit history length and employment
plt.figure(figsize=(10, 6))

# Create scatter plot with color by loan status
scatter = plt.scatter(
    df["cb_person_cred_hist_length"],
    df["person_emp_length"],
    c=df["loan_status"],
    alpha=0.5,
    cmap="coolwarm",
)

# Add regression line
z = np.polyfit(df["cb_person_cred_hist_length"], df["person_emp_length"],
p = np.poly1d(z)
plt.plot(
    df["cb_person_cred_hist_length"],
    p(df["cb_person_cred_hist_length"]),
    "k--",
    linewidth=2,
    label=f"Trend line: y={z[0]:.3f}x+{z[1]:.3f}",
)

plt.colorbar(scatter, label="Loan Status (1=Default)")
plt.xlabel("Credit History Length (Years)")
plt.ylabel("Employment Length (Years)")
plt.title("Relationship Between Credit History Length and Employment Leng
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



```
In [46]: credit_longer_than_employment = df[
    df["cb_person_cred_hist_length"] > df["person_emp_length"]
].copy()

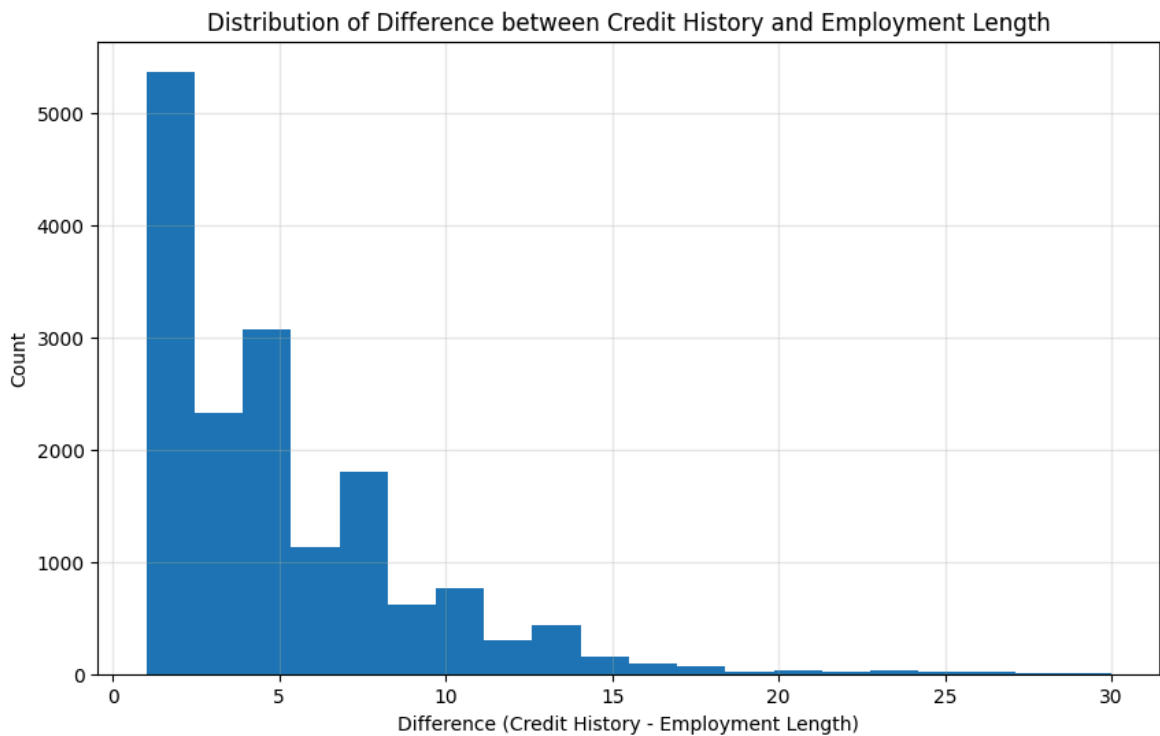
percentage = (len(credit_longer_than_employment) / len(df)) * 100

print(
    f"Number of people with credit history longer than employment: {len(c
)

credit_longer_than_employment["history_employment_diff"] = (
    credit_longer_than_employment["cb_person_cred_hist_length"]
    - credit_longer_than_employment["person_emp_length"]
)

plt.figure(figsize=(10, 6))
plt.hist(credit_longer_than_employment["history_employment_diff"], bins=2
plt.title("Distribution of Difference between Credit History and Employme
plt.xlabel("Difference (Credit History - Employment Length)")
plt.ylabel("Count")
plt.grid(True, alpha=0.3)
plt.show()
```

Number of people with credit history longer than employment: 16333 (51.81% of total)



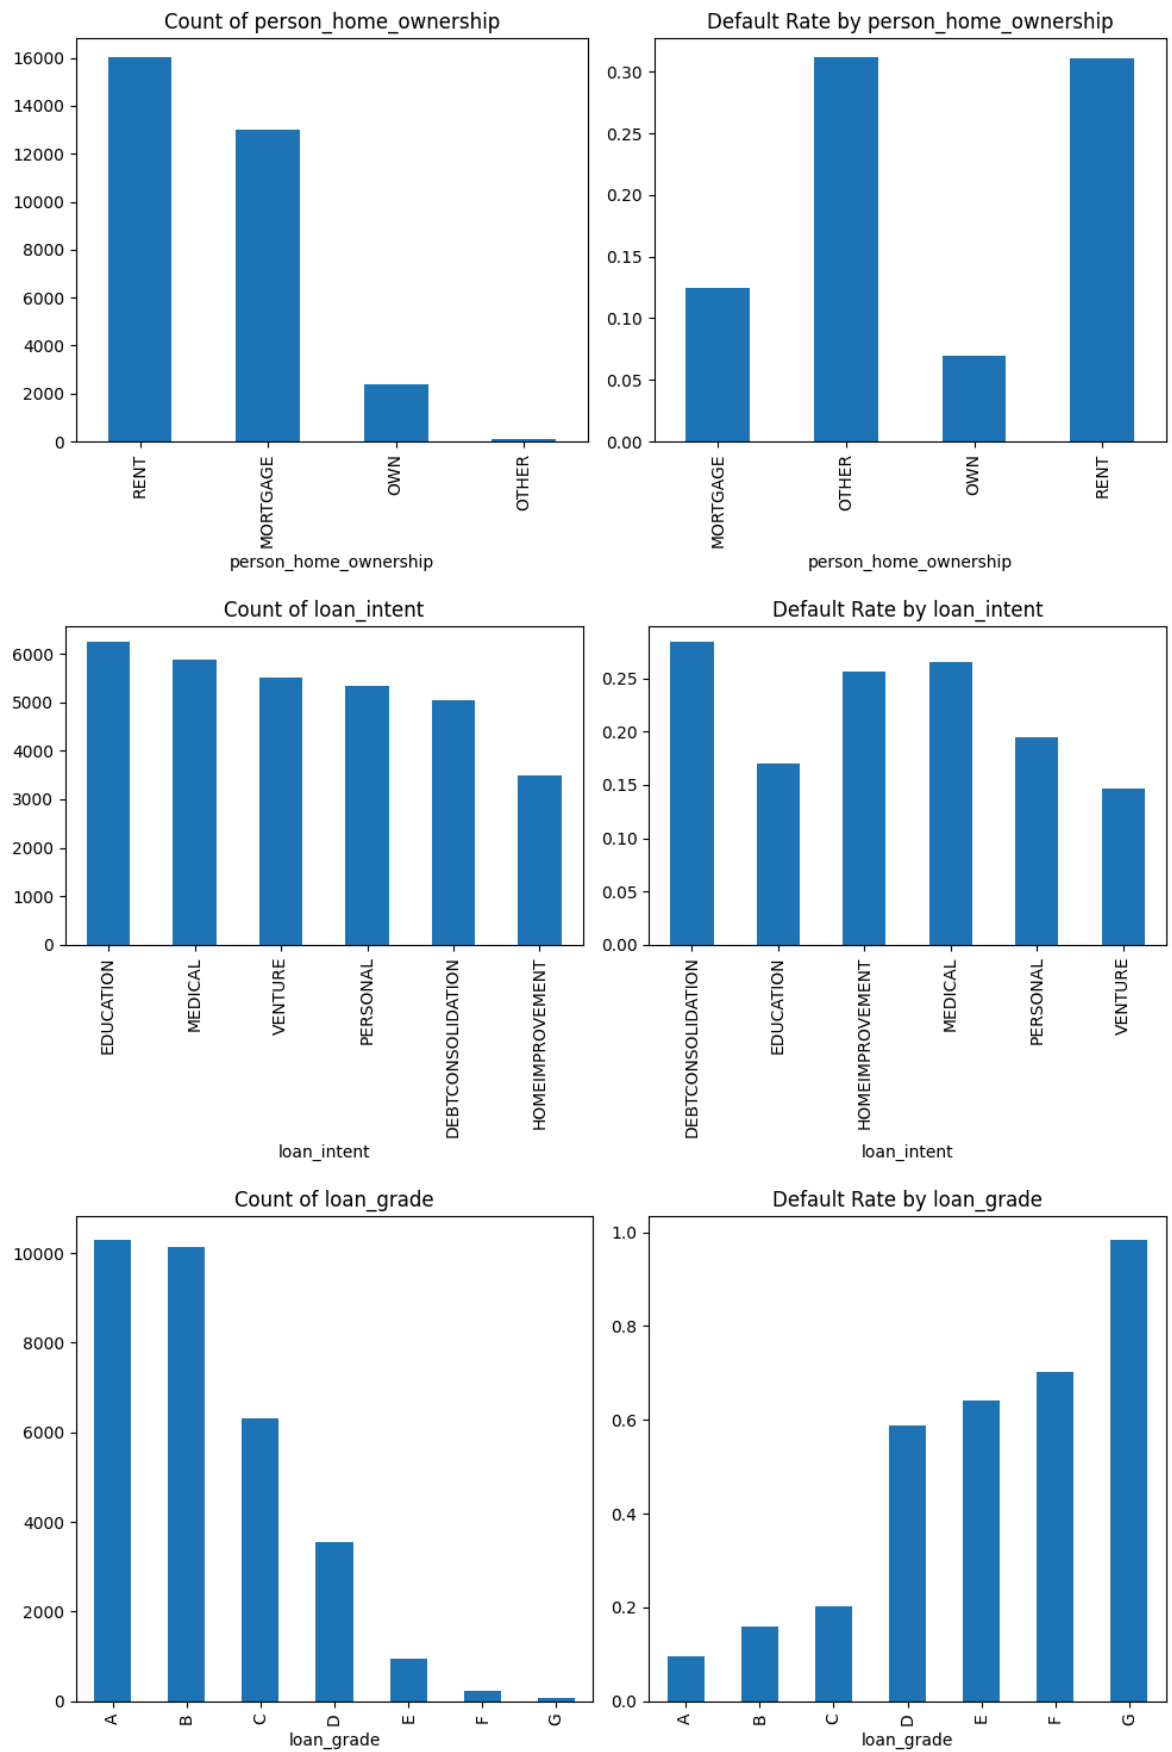
As you can see - too much of our data exhibits this peculiar behavior. Either there's some part of how credit history works in the US that we're not understanding, or that this synthetic dataset simply didn't model this relationship enough during generation. Either way, we decided to leave it as is.

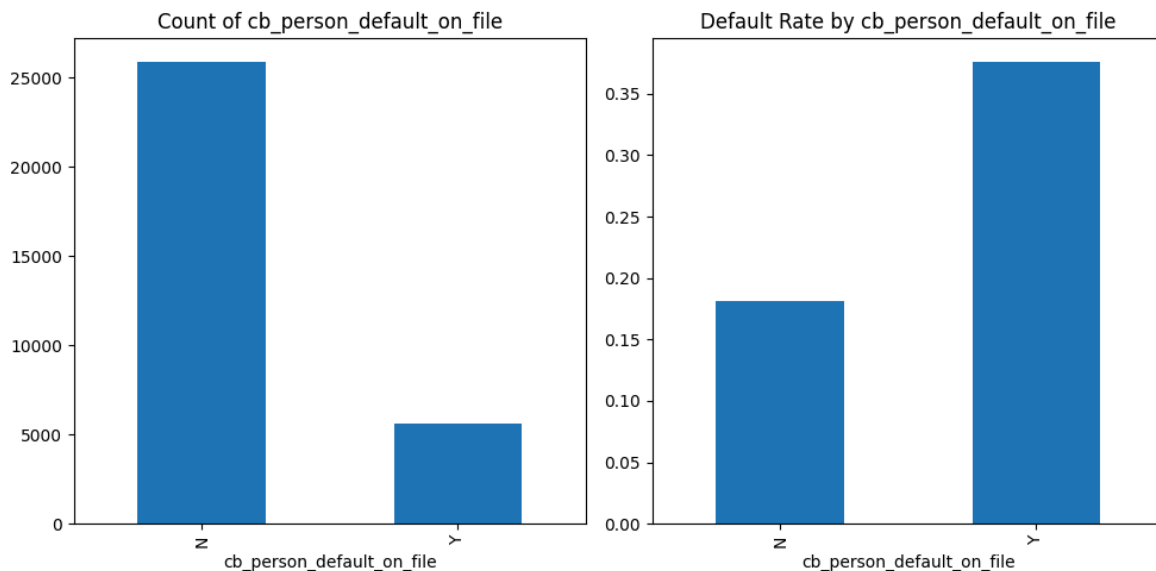
```
In [47]: # categorical features visualization
cat_features = [
    "person_home_ownership",
    "loan_intent",
    "loan_grade",
    "cb_person_default_on_file",
]

for feature in cat_features:
    plt.figure(figsize=(10, 5))

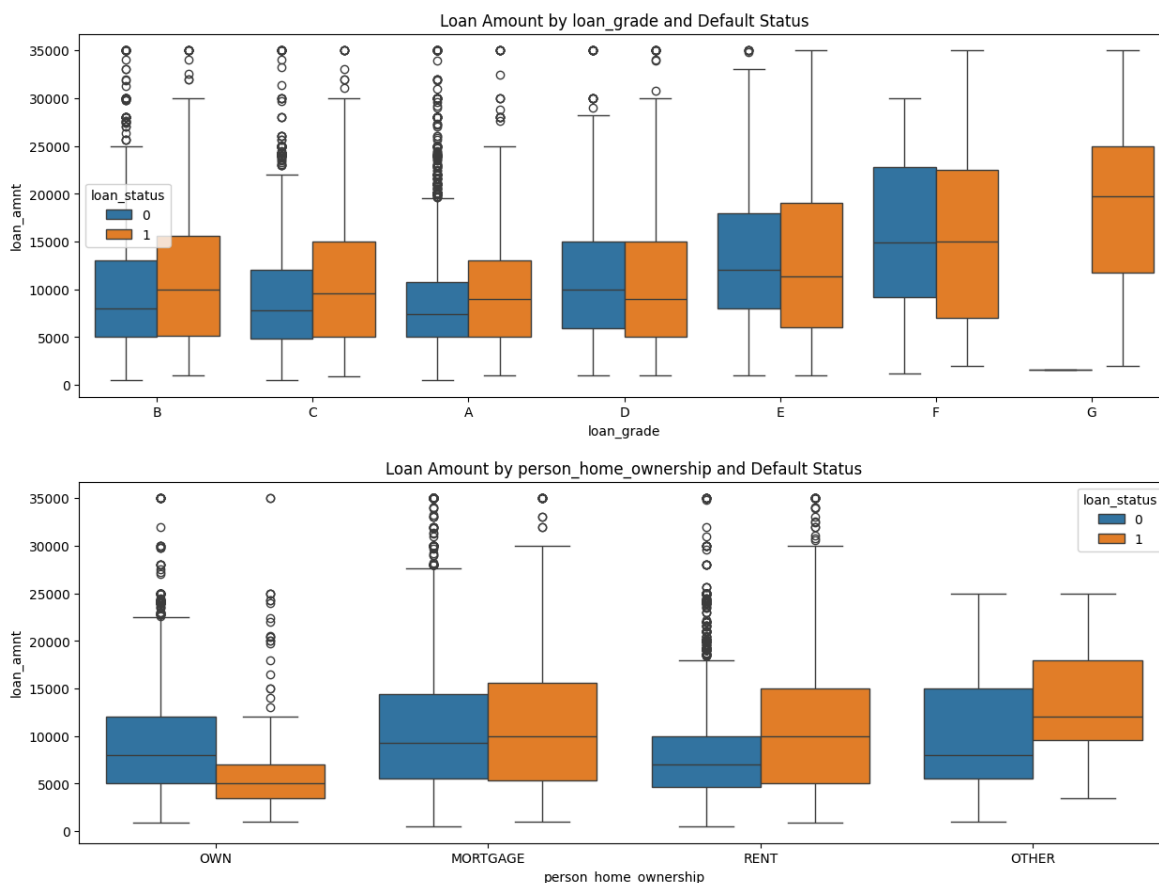
    # Count by category
    plt.subplot(1, 2, 1)
    df[feature].value_counts().plot(kind="bar")
    plt.title(f"Count of {feature}")

    # Default rate by category
    plt.subplot(1, 2, 2)
    df.groupby(feature)["loan_status"].mean().plot(kind="bar")
    plt.title(f"Default Rate by {feature}")
    plt.tight_layout()
    plt.show()
```





```
In [48]: # For key categorical variables, analyze numerical features by category
for cat in ["loan_grade", "person_home_ownership"]:
    plt.figure(figsize=(15, 5))
    sns.boxplot(x=cat, y="loan_amnt", hue="loan_status", data=df)
    plt.title(f"Loan Amount by {cat} and Default Status")
    plt.show()
```



## Data Preprocessing

Principally, categorical labels need to be encoded into numerical ones.

We also investigate two strategies --> impute missing data VS delete missing data records

One thing that stood out was how home ownership type = OTHER had a high 30% default rate. There are a few strategies:

1. Treat as RENT on the basis that their default rates are almost the same
2. Delete the record entirely
3. Do nothing

```
In [49]: from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split

# https://contrib.scikit-learn.org/category_encoders/
# other encoders we can try:
import category_encoders as ce
```

```
In [50]: encoder = ce.OneHotEncoder(
    cols=cat_features,
    return_df=True,
    use_cat_names=True,
)

df = encoder.fit_transform(df)
df
```

```
Out [50]:
```

	person_age	person_income	person_home_ownership_OWN	person_home_
1	21	9600	1	
2	25	9600	0	
3	23	65500	0	
4	24	54400	0	
5	21	9900	1	
...	...	...	...	...
32576	57	53000	0	
32577	54	120000	0	
32578	65	76000	0	
32579	56	150000	0	
32580	66	42000	0	

31527 rows × 27 columns

## Approach 1: Data Imputation

Based on checking the null values from above, the data that needs to be imputed are all numerical (person\_emp\_length and interest\_rate).

```
In [51]: X = df.drop("loan_status", axis=1)
y = df["loan_status"]
```

```
In [52]: # train test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [53]: # Standardize numerical features with StandardScaler
# Scale numerical features
scaler = StandardScaler()
numerical_cols = [
    col
    for col in X_train.columns
    if X_train[col].dtype in ["int64", "float64"]
    and "person_home_ownership_" != col # these are all categorical
    and "loan_intent_" != col
    and "loan_grade_" != col
    and "cb_person_default_on_file_" != col
]
```

```
In [54]: # Get the names of columns with missing values
columns_with_nulls = X_train.columns[X_train.isnull().any()].tolist()

# Impute missing values for numerical features
numerical_imputer = SimpleImputer(strategy="mean")

# Create a copy to avoid modifying the original
X_train_imputed = X_train.copy()
X_test_imputed = X_test.copy()

# Impute missing values
X_train_imputed[columns_with_nulls] = numerical_imputer.fit_transform(
    X_train[columns_with_nulls]
)
X_test_imputed[columns_with_nulls] = numerical_imputer.transform(
    X_test[columns_with_nulls]
)

X_train_imputed[numerical_cols] = scaler.fit_transform(X_train_imputed[numerical_cols])
X_test_imputed[numerical_cols] = scaler.transform(X_test_imputed[numerical_cols])

print(f"Training data shape: {X_train_imputed.shape}")
print(f"Testing data shape: {X_test_imputed.shape}")
print(f"Number of features: {X_train_imputed.shape[1]}")
```

Training data shape: (25221, 26)

Testing data shape: (6306, 26)

Number of features: 26

## Approach 2 (Our Actual Approach):

Doing data imputation requires us to split the data first (so that information of the mean doesn't leak across sets).

However, a different approach is to just drop the rows where there are null values present.

To keep things simple, we opted to continue the notebook by using the deletion method over imputation. (We did not use the data imputed above). We thus continue our analysis below using this simplistic method.

## Introducing Pycaret

PyCaret is a framework that allows us to test the dataset against different kinds of classifiers out there. It is considered an auto-ML framework in that once we have our dataset prepared properly, it basically does everything else for us.

```
In [55]: from pycaret.classification import *  
  
model = setup(data=df, target="loan_status")
```



	Description	Value
0	Session id	8770
1	Target	loan_status
2	Target type	Binary
3	Original data shape	(31527, 27)
4	Transformed data shape	(31527, 27)
5	Transformed train set shape	(22068, 27)
6	Transformed test set shape	(9459, 27)
7	Numeric features	26
8	Rows with missing values	9.6%
9	Preprocess	True
10	Imputation type	simple
11	Numeric imputation	mean
12	Categorical imputation	mode
13	Fold Generator	StratifiedKFold
14	Fold Number	10
15	CPU Jobs	-1
16	Use GPU	False
17	Log Experiment	False
18	Experiment Name	clf-default-name
19	USI	e6d6

PyCaret is an AutoML framework that uses a collection of different classifiers to see which performs the best. We use this as a reference for what we should expect to get with our own implementations.

```
In [56]: best = compare_models()
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	T
<b>lightgbm</b>	Light Gradient Boosting Machine	0.9371	0.9468	0.7274	0.9750	0.8329	0.7952	0.8083	0
<b>xgboost</b>	Extreme Gradient Boosting	0.9366	0.9489	0.7434	0.9528	0.8348	0.7964	0.8060	0
<b>rf</b>	Random Forest Classifier	0.9346	0.9302	0.7198	0.9696	0.8259	0.7868	0.8002	0
<b>gbc</b>	Gradient Boosting Classifier	0.9280	0.9278	0.7110	0.9415	0.8098	0.7665	0.7782	0
<b>et</b>	Extra Trees Classifier	0.9205	0.9134	0.6861	0.9265	0.7881	0.7405	0.7533	0
<b>dt</b>	Decision Tree Classifier	0.8893	0.8446	0.7660	0.7331	0.7490	0.6780	0.6784	0
<b>ada</b>	Ada Boost Classifier	0.8866	0.8989	0.6447	0.7916	0.7104	0.6408	0.6462	0
<b>lda</b>	Linear Discriminant Analysis	0.8655	0.8689	0.5883	0.7361	0.6537	0.5715	0.5772	0
<b>ridge</b>	Ridge Classifier	0.8610	0.8688	0.5190	0.7614	0.6169	0.5358	0.5506	0
<b>lr</b>	Logistic Regression	0.8537	0.8563	0.4701	0.7611	0.5810	0.4983	0.5197	0
<b>knn</b>	K Neighbors Classifier	0.8317	0.8087	0.4932	0.6442	0.5583	0.4567	0.4631	0
<b>nb</b>	Naive Bayes	0.8157	0.7675	0.3244	0.6480	0.4311	0.3356	0.3648	0
<b>qda</b>	Quadratic Discriminant Analysis	0.7964	0.8644	0.1234	0.6523	0.2072	0.1488	0.2185	0
<b>dummy</b>	Dummy Classifier	0.7841	0.5000	0.0000	0.0000	0.0000	0.0000	0.0000	0
<b>svm</b>	SVM - Linear Kernel	0.7557	0.7335	0.2249	0.4004	0.2031	0.1322	0.1645	0

```
In [57]: evaluate_model(best)
```

```
interactive(children=(ToggleButtons(description='Plot Type:', icons=('',)),
options= (('Pipeline Plot', 'pipelin...
```

```
In [58]: from sklearn.decomposition import PCA

scaler = StandardScaler()
```

```

df_clean = df.dropna()
scaled_data = scaler.fit_transform(df_clean)

# Perform PCA
pca = PCA(n_components=2)
principal_components = pca.fit_transform(scaled_data)

# Convert to DataFrame
pca_df = pd.DataFrame(data=principal_components, columns=["PC1", "PC2"])

# Show explained variance
print(pca.explained_variance_ratio_)

```

```
[0.13189015 0.08532481]
```

In [59]: `print(pca_df)`

```

          PC1      PC2
0    -0.957522 -1.430459
1     1.319681 -0.939564
2     2.541408 -1.666157
3     4.579694 -0.319884
4    -1.133638 -1.694136
...         ...     ...
28494 -0.046952  4.685323
28495 -2.069325  4.001129
28496  0.890104  2.914141
28497 -1.179385  4.666782
28498 -0.598310  2.517073

```

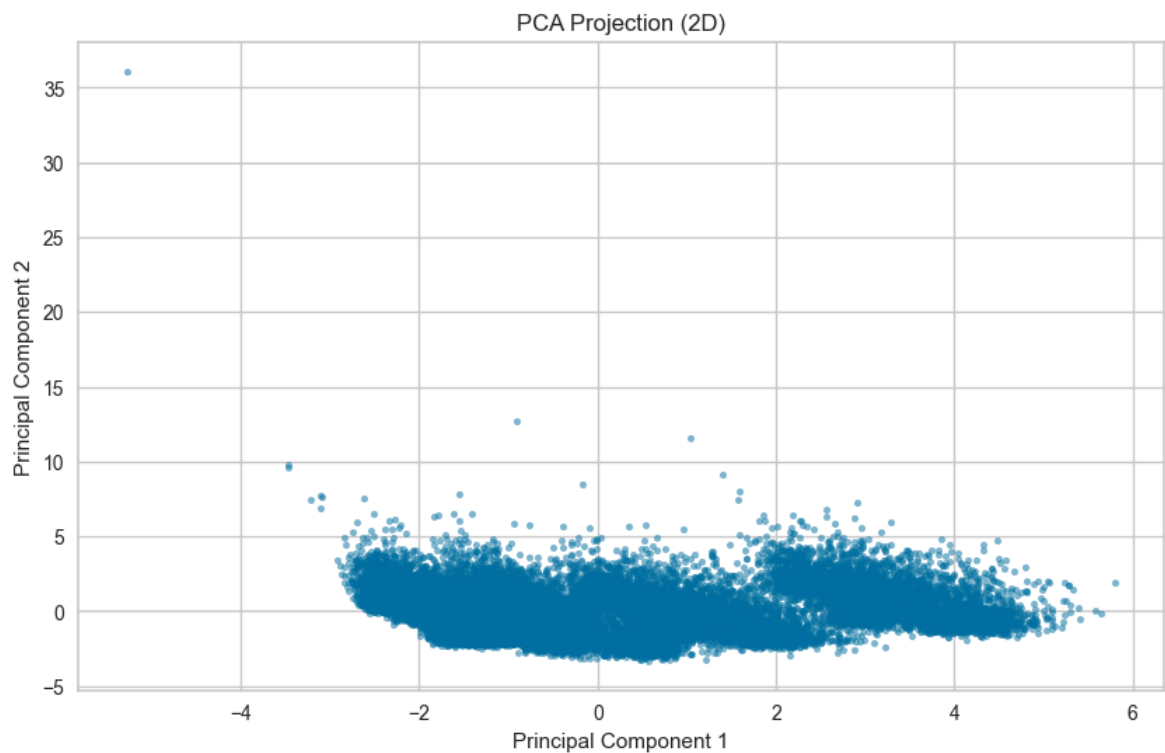
```
[28499 rows x 2 columns]
```

In [60]:

```

plt.figure(figsize=(10, 6))
plt.scatter(pca_df["PC1"], pca_df["PC2"], alpha=0.5, s=10)
plt.title("PCA Projection (2D)")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.grid(True)
plt.show()

```



```
In [61]: df_clean = df.dropna()
scaler = StandardScaler()
scaled_data = scaler.fit_transform(df_clean)

pca = PCA(n_components=3)
principal_components = pca.fit_transform(scaled_data)

pca_df = pd.DataFrame(principal_components, columns=["PC1", "PC2", "PC3"])

# Show explained variance
print(pca.explained_variance_ratio_)
```

```
[0.13188918 0.08530775 0.07152489]
```

Low explained variance for each Principal Components suggests that most of the important variation is spread out across more components (data is likely high-dimensional and complex), meaning we have to keep more features for a good model. Also means that linear models will likely struggle.


```
In [62]: fig = plt.figure(figsize=(10, 7))
ax = fig.add_subplot(111, projection="3d")

ax.scatter(pca_df["PC1"], pca_df["PC2"], pca_df["PC3"], alpha=0.6, s=10)

ax.set_title("3D PCA Projection")
ax.set_xlabel("Principal Component 1")
ax.set_ylabel("Principal Component 2")
ax.set_zlabel("Principal Component 3")

plt.show()
```

### 3D PCA Projection



Visually it appears that there are 3 distinct clusters

Since data is highly complex, try Neural Network which can effectively capture non linear relationships

```
In [63]: import torch
import torch.nn as nn
import torch.optim as optim
```

```
In [33]: df_cleaned = df.dropna()
X = df_cleaned.drop("loan_status", axis=1).to_numpy(dtype=np.float32)
y = df_cleaned["loan_status"].to_numpy(dtype=np.float32).reshape(-1, 1)

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

# Convert to tensors
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
```

```
y_train_tensor = torch.tensor(y_train, dtype=torch.float32)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test, dtype=torch.float32)
```

```
In [64]: assert not torch.isnan(X_train_tensor).any(), "NaNs in training data"
assert not torch.isnan(y_train_tensor).any(), "NaNs in training labels"
```

Below are our trial and error attempts at modifying the neural network to produce better results. We produce models that demonstrate about 92-93% accuracy on the held out test set.

```
In [65]: class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.fc1 = nn.Linear(input_size, 32)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(32, 16)
        self.relu2 = nn.ReLU()
        self.fc3 = nn.Linear(16, 1)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = self.relu1(self.fc1(x))
        x = self.relu2(self.fc2(x))
        x = self.sigmoid(self.fc3(x))
        return x

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training loop
epochs = 1000
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

# Evaluation
model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/1000], Loss: 0.7343  
Epoch [20/1000], Loss: 0.7040  
Epoch [30/1000], Loss: 0.6732  
Epoch [40/1000], Loss: 0.6385  
Epoch [50/1000], Loss: 0.5990  
Epoch [60/1000], Loss: 0.5561  
Epoch [70/1000], Loss: 0.5131  
Epoch [80/1000], Loss: 0.4739  
Epoch [90/1000], Loss: 0.4405  
Epoch [100/1000], Loss: 0.4127  
Epoch [110/1000], Loss: 0.3897  
Epoch [120/1000], Loss: 0.3708  
Epoch [130/1000], Loss: 0.3552  
Epoch [140/1000], Loss: 0.3423  
Epoch [150/1000], Loss: 0.3317  
Epoch [160/1000], Loss: 0.3229  
Epoch [170/1000], Loss: 0.3156  
Epoch [180/1000], Loss: 0.3093  
Epoch [190/1000], Loss: 0.3039  
Epoch [200/1000], Loss: 0.2989  
Epoch [210/1000], Loss: 0.2945  
Epoch [220/1000], Loss: 0.2904  
Epoch [230/1000], Loss: 0.2867  
Epoch [240/1000], Loss: 0.2833  
Epoch [250/1000], Loss: 0.2803  
Epoch [260/1000], Loss: 0.2775  
Epoch [270/1000], Loss: 0.2749  
Epoch [280/1000], Loss: 0.2725  
Epoch [290/1000], Loss: 0.2702  
Epoch [300/1000], Loss: 0.2680  
Epoch [310/1000], Loss: 0.2660  
Epoch [320/1000], Loss: 0.2639  
Epoch [330/1000], Loss: 0.2619  
Epoch [340/1000], Loss: 0.2600  
Epoch [350/1000], Loss: 0.2582  
Epoch [360/1000], Loss: 0.2564  
Epoch [370/1000], Loss: 0.2547  
Epoch [380/1000], Loss: 0.2531  
Epoch [390/1000], Loss: 0.2516  
Epoch [400/1000], Loss: 0.2501  
Epoch [410/1000], Loss: 0.2486  
Epoch [420/1000], Loss: 0.2472  
Epoch [430/1000], Loss: 0.2458  
Epoch [440/1000], Loss: 0.2445  
Epoch [450/1000], Loss: 0.2433  
Epoch [460/1000], Loss: 0.2421  
Epoch [470/1000], Loss: 0.2410  
Epoch [480/1000], Loss: 0.2400  
Epoch [490/1000], Loss: 0.2390  
Epoch [500/1000], Loss: 0.2381  
Epoch [510/1000], Loss: 0.2372  
Epoch [520/1000], Loss: 0.2364  
Epoch [530/1000], Loss: 0.2356  
Epoch [540/1000], Loss: 0.2349  
Epoch [550/1000], Loss: 0.2341  
Epoch [560/1000], Loss: 0.2334  
Epoch [570/1000], Loss: 0.2327  
Epoch [580/1000], Loss: 0.2320  
Epoch [590/1000], Loss: 0.2313  
Epoch [600/1000], Loss: 0.2306

```
Epoch [610/1000], Loss: 0.2300
Epoch [620/1000], Loss: 0.2293
Epoch [630/1000], Loss: 0.2287
Epoch [640/1000], Loss: 0.2281
Epoch [650/1000], Loss: 0.2275
Epoch [660/1000], Loss: 0.2269
Epoch [670/1000], Loss: 0.2263
Epoch [680/1000], Loss: 0.2257
Epoch [690/1000], Loss: 0.2251
Epoch [700/1000], Loss: 0.2245
Epoch [710/1000], Loss: 0.2239
Epoch [720/1000], Loss: 0.2233
Epoch [730/1000], Loss: 0.2228
Epoch [740/1000], Loss: 0.2223
Epoch [750/1000], Loss: 0.2218
Epoch [760/1000], Loss: 0.2212
Epoch [770/1000], Loss: 0.2207
Epoch [780/1000], Loss: 0.2202
Epoch [790/1000], Loss: 0.2196
Epoch [800/1000], Loss: 0.2191
Epoch [810/1000], Loss: 0.2186
Epoch [820/1000], Loss: 0.2180
Epoch [830/1000], Loss: 0.2175
Epoch [840/1000], Loss: 0.2170
Epoch [850/1000], Loss: 0.2166
Epoch [860/1000], Loss: 0.2162
Epoch [870/1000], Loss: 0.2157
Epoch [880/1000], Loss: 0.2153
Epoch [890/1000], Loss: 0.2149
Epoch [900/1000], Loss: 0.2145
Epoch [910/1000], Loss: 0.2141
Epoch [920/1000], Loss: 0.2137
Epoch [930/1000], Loss: 0.2133
Epoch [940/1000], Loss: 0.2129
Epoch [950/1000], Loss: 0.2125
Epoch [960/1000], Loss: 0.2121
Epoch [970/1000], Loss: 0.2117
Epoch [980/1000], Loss: 0.2114
Epoch [990/1000], Loss: 0.2110
Epoch [1000/1000], Loss: 0.2107
```

Train Accuracy: 0.9308

Test Accuracy: 0.9174

Large difference between train and test accuracy implies overfitting

Add dropout to reduce overfitting

In [67]: `first_layer_neurons = 128`

```
# Increase the number of neurons in each layer by a factor of 4
class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.fc1 = nn.Linear(input_size, first_layer_neurons)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(first_layer_neurons, first_layer_neurons // 2)
        self.relu2 = nn.ReLU()
```



```
self.fc3 = nn.Linear(first_layer_neurons // 2, 1)
self.sigmoid = nn.Sigmoid()
self.dropout1 = nn.Dropout(0.3)
self.dropout2 = nn.Dropout(0.3)

def forward(self, x):
    x = self.relu1(self.fc1(x))
    x = self.dropout1(x)
    x = self.relu2(self.fc2(x))
    x = self.dropout2(x)
    x = self.sigmoid(self.fc3(x))
    return x

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training loop
epochs = 1000
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/1000], Loss: 0.5418  
Epoch [20/1000], Loss: 0.4720  
Epoch [30/1000], Loss: 0.4164  
Epoch [40/1000], Loss: 0.3817  
Epoch [50/1000], Loss: 0.3580  
Epoch [60/1000], Loss: 0.3433  
Epoch [70/1000], Loss: 0.3306  
Epoch [80/1000], Loss: 0.3196  
Epoch [90/1000], Loss: 0.3132  
Epoch [100/1000], Loss: 0.3054  
Epoch [110/1000], Loss: 0.3020  
Epoch [120/1000], Loss: 0.2948  
Epoch [130/1000], Loss: 0.2926  
Epoch [140/1000], Loss: 0.2873  
Epoch [150/1000], Loss: 0.2851  
Epoch [160/1000], Loss: 0.2810  
Epoch [170/1000], Loss: 0.2807  
Epoch [180/1000], Loss: 0.2787  
Epoch [190/1000], Loss: 0.2717  
Epoch [200/1000], Loss: 0.2703  
Epoch [210/1000], Loss: 0.2680  
Epoch [220/1000], Loss: 0.2659  
Epoch [230/1000], Loss: 0.2616  
Epoch [240/1000], Loss: 0.2607  
Epoch [250/1000], Loss: 0.2596  
Epoch [260/1000], Loss: 0.2572  
Epoch [270/1000], Loss: 0.2560  
Epoch [280/1000], Loss: 0.2547  
Epoch [290/1000], Loss: 0.2538  
Epoch [300/1000], Loss: 0.2521  
Epoch [310/1000], Loss: 0.2489  
Epoch [320/1000], Loss: 0.2507  
Epoch [330/1000], Loss: 0.2489  
Epoch [340/1000], Loss: 0.2460  
Epoch [350/1000], Loss: 0.2438  
Epoch [360/1000], Loss: 0.2418  
Epoch [370/1000], Loss: 0.2435  
Epoch [380/1000], Loss: 0.2413  
Epoch [390/1000], Loss: 0.2399  
Epoch [400/1000], Loss: 0.2393  
Epoch [410/1000], Loss: 0.2374  
Epoch [420/1000], Loss: 0.2369  
Epoch [430/1000], Loss: 0.2357  
Epoch [440/1000], Loss: 0.2372  
Epoch [450/1000], Loss: 0.2342  
Epoch [460/1000], Loss: 0.2324  
Epoch [470/1000], Loss: 0.2328  
Epoch [480/1000], Loss: 0.2328  
Epoch [490/1000], Loss: 0.2305  
Epoch [500/1000], Loss: 0.2325  
Epoch [510/1000], Loss: 0.2303  
Epoch [520/1000], Loss: 0.2267  
Epoch [530/1000], Loss: 0.2292  
Epoch [540/1000], Loss: 0.2273  
Epoch [550/1000], Loss: 0.2261  
Epoch [560/1000], Loss: 0.2265  
Epoch [570/1000], Loss: 0.2233  
Epoch [580/1000], Loss: 0.2256  
Epoch [590/1000], Loss: 0.2230  
Epoch [600/1000], Loss: 0.2232

```
Epoch [610/1000], Loss: 0.2232
Epoch [620/1000], Loss: 0.2234
Epoch [630/1000], Loss: 0.2209
Epoch [640/1000], Loss: 0.2205
Epoch [650/1000], Loss: 0.2225
Epoch [660/1000], Loss: 0.2210
Epoch [670/1000], Loss: 0.2202
Epoch [680/1000], Loss: 0.2200
Epoch [690/1000], Loss: 0.2186
Epoch [700/1000], Loss: 0.2177
Epoch [710/1000], Loss: 0.2177
Epoch [720/1000], Loss: 0.2196
Epoch [730/1000], Loss: 0.2174
Epoch [740/1000], Loss: 0.2167
Epoch [750/1000], Loss: 0.2173
Epoch [760/1000], Loss: 0.2153
Epoch [770/1000], Loss: 0.2161
Epoch [780/1000], Loss: 0.2184
Epoch [790/1000], Loss: 0.2162
Epoch [800/1000], Loss: 0.2150
Epoch [810/1000], Loss: 0.2148
Epoch [820/1000], Loss: 0.2123
Epoch [830/1000], Loss: 0.2133
Epoch [840/1000], Loss: 0.2147
Epoch [850/1000], Loss: 0.2120
Epoch [860/1000], Loss: 0.2119
Epoch [870/1000], Loss: 0.2130
Epoch [880/1000], Loss: 0.2118
Epoch [890/1000], Loss: 0.2114
Epoch [900/1000], Loss: 0.2132
Epoch [910/1000], Loss: 0.2109
Epoch [920/1000], Loss: 0.2131
Epoch [930/1000], Loss: 0.2133
Epoch [940/1000], Loss: 0.2116
Epoch [950/1000], Loss: 0.2113
Epoch [960/1000], Loss: 0.2108
Epoch [970/1000], Loss: 0.2093
Epoch [980/1000], Loss: 0.2081
Epoch [990/1000], Loss: 0.2095
Epoch [1000/1000], Loss: 0.2100
```

Train Accuracy: 0.9336

Test Accuracy: 0.9261

As expected, test accuracy increased

We can try to improve the overall accuracy by adding more layers

Try lesser epochs to reduce overfitting

```
In [68]: first_layer_neurons = 1024

# Increase the number of neurons in each layer by a factor of 4
class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.fc1 = nn.Linear(input_size, first_layer_neurons)
```

```

self.relu1 = nn.ReLU()
self.fc2 = nn.Linear(first_layer_neurons, first_layer_neurons //
self.relu2 = nn.ReLU()
self.fc3 = nn.Linear(first_layer_neurons // 2, first_layer_neurons
self.relu3 = nn.ReLU()
self.fc4 = nn.Linear(first_layer_neurons // 4, first_layer_neurons
self.relu4 = nn.ReLU()
self.fc5 = nn.Linear(first_layer_neurons // 8, 1)
self.relu5 = nn.ReLU()
self.sigmoid = nn.Sigmoid()
self.dropout1 = nn.Dropout(0.3)
self.dropout2 = nn.Dropout(0.3)
self.dropout3 = nn.Dropout(0.3)
self.dropout4 = nn.Dropout(0.3)
self.dropout5 = nn.Dropout(0.3)

def forward(self, x):
    x = self.relu1(self.fc1(x))
    x = self.dropout1(x)
    x = self.relu2(self.fc2(x))
    x = self.dropout2(x)
    x = self.relu3(self.fc3(x))
    x = self.dropout3(x)
    x = self.relu4(self.fc4(x))
    x = self.dropout4(x)
    x = self.fc5(x)
    x = self.sigmoid(x)
    return x

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training loop
epochs = 200
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().me
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")

```

```
Epoch [10/200], Loss: 0.4419
Epoch [20/200], Loss: 0.3765
Epoch [30/200], Loss: 0.3497
Epoch [40/200], Loss: 0.3111
Epoch [50/200], Loss: 0.2821
Epoch [60/200], Loss: 0.2672
Epoch [70/200], Loss: 0.2559
Epoch [80/200], Loss: 0.2486
Epoch [90/200], Loss: 0.2402
Epoch [100/200], Loss: 0.2354
Epoch [110/200], Loss: 0.2279
Epoch [120/200], Loss: 0.2242
Epoch [130/200], Loss: 0.2212
Epoch [140/200], Loss: 0.2185
Epoch [150/200], Loss: 0.2163
Epoch [160/200], Loss: 0.2135
Epoch [170/200], Loss: 0.2103
Epoch [180/200], Loss: 0.2071
Epoch [190/200], Loss: 0.2045
Epoch [200/200], Loss: 0.2016
```

Train Accuracy: 0.9377

Test Accuracy: 0.9261

Change from sigmoid to BCEWithLogitsLoss which is more stable by avoiding issues with vanishing gradients

In [69]: first\_layer\_neurons = 1024

```
# Increase the number of neurons in each layer by a factor of 4
class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.fc1 = nn.Linear(input_size, first_layer_neurons)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(first_layer_neurons, first_layer_neurons // 2)
        self.relu2 = nn.ReLU()
        self.fc3 = nn.Linear(first_layer_neurons // 2, first_layer_neurons // 4)
        self.relu3 = nn.ReLU()
        self.fc4 = nn.Linear(first_layer_neurons // 4, first_layer_neurons // 8)
        self.relu4 = nn.ReLU()
        self.fc5 = nn.Linear(first_layer_neurons // 8, 1)
        self.relu5 = nn.ReLU()
        self.dropout1 = nn.Dropout(0.3)
        self.dropout2 = nn.Dropout(0.3)
        self.dropout3 = nn.Dropout(0.3)
        self.dropout4 = nn.Dropout(0.3)
        self.dropout5 = nn.Dropout(0.3)

    def forward(self, x):
        x = self.relu1(self.fc1(x))
        x = self.dropout1(x)
        x = self.relu2(self.fc2(x))
        x = self.dropout2(x)
        x = self.relu3(self.fc3(x))
        x = self.dropout3(x)
        x = self.relu4(self.fc4(x))
```

```

        x = self.dropout4(x)
        x = self.fc5(x)
        return x

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training loop
epochs = 200
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")

```

```

Epoch [10/200], Loss: 0.4184
Epoch [20/200], Loss: 0.3683
Epoch [30/200], Loss: 0.3243
Epoch [40/200], Loss: 0.2886
Epoch [50/200], Loss: 0.2697
Epoch [60/200], Loss: 0.2587
Epoch [70/200], Loss: 0.2490
Epoch [80/200], Loss: 0.2401
Epoch [90/200], Loss: 0.2343
Epoch [100/200], Loss: 0.2290
Epoch [110/200], Loss: 0.2245
Epoch [120/200], Loss: 0.2199
Epoch [130/200], Loss: 0.2194
Epoch [140/200], Loss: 0.2144
Epoch [150/200], Loss: 0.2093
Epoch [160/200], Loss: 0.2088
Epoch [170/200], Loss: 0.2046
Epoch [180/200], Loss: 0.2043
Epoch [190/200], Loss: 0.2024
Epoch [200/200], Loss: 0.1972

```

Train Accuracy: 0.9375

Test Accuracy: 0.9219

Increase epoch to reduce bias and add weight decay (L2 regularization) to reduce variance/overfitting

```
In [70]: first_layer_neurons = 1024

# Increase the number of neurons in each layer by a factor of 4
class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.fc1 = nn.Linear(input_size, first_layer_neurons)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(first_layer_neurons, first_layer_neurons // 2)
        self.relu2 = nn.ReLU()
        self.fc3 = nn.Linear(first_layer_neurons // 2, first_layer_neurons // 4)
        self.relu3 = nn.ReLU()
        self.fc4 = nn.Linear(first_layer_neurons // 4, first_layer_neurons // 8)
        self.relu4 = nn.ReLU()
        self.fc5 = nn.Linear(first_layer_neurons // 8, 1)
        self.relu5 = nn.ReLU()
        self.dropout1 = nn.Dropout(0.3)
        self.dropout2 = nn.Dropout(0.3)
        self.dropout3 = nn.Dropout(0.3)
        self.dropout4 = nn.Dropout(0.3)
        self.dropout5 = nn.Dropout(0.3)

    def forward(self, x):
        x = self.relu1(self.fc1(x))
        x = self.dropout1(x)
        x = self.relu2(self.fc2(x))
        x = self.dropout2(x)
        x = self.relu3(self.fc3(x))
        x = self.dropout3(x)
        x = self.relu4(self.fc4(x))
        x = self.dropout4(x)
        x = self.fc5(x)
        return x

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-5)

# Training loop
epochs = 500
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
```

```
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().me
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```



Epoch [10/500], Loss: 0.4348  
Epoch [20/500], Loss: 0.3690  
Epoch [30/500], Loss: 0.3272  
Epoch [40/500], Loss: 0.2910  
Epoch [50/500], Loss: 0.2707  
Epoch [60/500], Loss: 0.2592  
Epoch [70/500], Loss: 0.2503  
Epoch [80/500], Loss: 0.2440  
Epoch [90/500], Loss: 0.2366  
Epoch [100/500], Loss: 0.2317  
Epoch [110/500], Loss: 0.2266  
Epoch [120/500], Loss: 0.2222  
Epoch [130/500], Loss: 0.2213  
Epoch [140/500], Loss: 0.2168  
Epoch [150/500], Loss: 0.2131  
Epoch [160/500], Loss: 0.2107  
Epoch [170/500], Loss: 0.2066  
Epoch [180/500], Loss: 0.2060  
Epoch [190/500], Loss: 0.2020  
Epoch [200/500], Loss: 0.1994  
Epoch [210/500], Loss: 0.1970  
Epoch [220/500], Loss: 0.1955  
Epoch [230/500], Loss: 0.1951  
Epoch [240/500], Loss: 0.1949  
Epoch [250/500], Loss: 0.1877  
Epoch [260/500], Loss: 0.1876  
Epoch [270/500], Loss: 0.1858  
Epoch [280/500], Loss: 0.1838  
Epoch [290/500], Loss: 0.1789  
Epoch [300/500], Loss: 0.1768  
Epoch [310/500], Loss: 0.1767  
Epoch [320/500], Loss: 0.1750  
Epoch [330/500], Loss: 0.1735  
Epoch [340/500], Loss: 0.1707  
Epoch [350/500], Loss: 0.1687  
Epoch [360/500], Loss: 0.1667  
Epoch [370/500], Loss: 0.1662  
Epoch [380/500], Loss: 0.1634  
Epoch [390/500], Loss: 0.1601  
Epoch [400/500], Loss: 0.1567  
Epoch [410/500], Loss: 0.1604  
Epoch [420/500], Loss: 0.1530  
Epoch [430/500], Loss: 0.1517  
Epoch [440/500], Loss: 0.1485  
Epoch [450/500], Loss: 0.1456  
Epoch [460/500], Loss: 0.1450  
Epoch [470/500], Loss: 0.1479  
Epoch [480/500], Loss: 0.1414  
Epoch [490/500], Loss: 0.1391  
Epoch [500/500], Loss: 0.1386

Train Accuracy: 0.9566

Test Accuracy: 0.9186

Try increase dropout to prevent overfitting

Increase weight decay by factor of 10 to reduce overfitting and increase dropout to 0.5

```

In [71]: first_layer_neurons = 1024
         dropout_rate = 0.4

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.fc1 = nn.Linear(input_size, first_layer_neurons)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(first_layer_neurons, first_layer_neurons // 2)
        self.relu2 = nn.ReLU()
        self.fc3 = nn.Linear(first_layer_neurons // 2, first_layer_neurons // 4)
        self.relu3 = nn.ReLU()
        self.fc4 = nn.Linear(first_layer_neurons // 4, first_layer_neurons // 8)
        self.relu4 = nn.ReLU()
        self.fc5 = nn.Linear(first_layer_neurons // 8, 1)
        self.relu5 = nn.ReLU()
        self.dropout1 = nn.Dropout(dropout_rate)
        self.dropout2 = nn.Dropout(dropout_rate)
        self.dropout3 = nn.Dropout(dropout_rate)
        self.dropout4 = nn.Dropout(dropout_rate)
        self.dropout5 = nn.Dropout(dropout_rate)

    def forward(self, x):
        x = self.relu1(self.fc1(x))
        x = self.dropout1(x)
        x = self.relu2(self.fc2(x))
        x = self.dropout2(x)
        x = self.relu3(self.fc3(x))
        x = self.dropout3(x)
        x = self.relu4(self.fc4(x))
        x = self.dropout4(x)
        x = self.fc5(x)
        return x

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 500
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()

```

```
test_accuracy = (test_predicted_classes == y_train_tensor).float().me
print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
preds = model(X_test_tensor)
predicted_classes = (preds > 0.5).float()
accuracy = (predicted_classes == y_test_tensor).float().mean()
print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

```
Epoch [10/500], Loss: 0.4170
Epoch [20/500], Loss: 0.3730
Epoch [30/500], Loss: 0.3425
Epoch [40/500], Loss: 0.3022
Epoch [50/500], Loss: 0.2809
Epoch [60/500], Loss: 0.2661
Epoch [70/500], Loss: 0.2596
Epoch [80/500], Loss: 0.2489
Epoch [90/500], Loss: 0.2429
Epoch [100/500], Loss: 0.2363
Epoch [110/500], Loss: 0.2324
Epoch [120/500], Loss: 0.2293
Epoch [130/500], Loss: 0.2278
Epoch [140/500], Loss: 0.2237
Epoch [150/500], Loss: 0.2215
Epoch [160/500], Loss: 0.2198
Epoch [170/500], Loss: 0.2181
Epoch [180/500], Loss: 0.2134
Epoch [190/500], Loss: 0.2169
Epoch [200/500], Loss: 0.2131
Epoch [210/500], Loss: 0.2120
Epoch [220/500], Loss: 0.2105
Epoch [230/500], Loss: 0.2096
Epoch [240/500], Loss: 0.2074
Epoch [250/500], Loss: 0.2055
Epoch [260/500], Loss: 0.2058
Epoch [270/500], Loss: 0.2037
Epoch [280/500], Loss: 0.2052
Epoch [290/500], Loss: 0.2030
Epoch [300/500], Loss: 0.2020
Epoch [310/500], Loss: 0.2007
Epoch [320/500], Loss: 0.2003
Epoch [330/500], Loss: 0.2031
Epoch [340/500], Loss: 0.1994
Epoch [350/500], Loss: 0.1984
Epoch [360/500], Loss: 0.1962
Epoch [370/500], Loss: 0.1949
Epoch [380/500], Loss: 0.1938
Epoch [390/500], Loss: 0.1938
Epoch [400/500], Loss: 0.1926
Epoch [410/500], Loss: 0.1913
Epoch [420/500], Loss: 0.1913
Epoch [430/500], Loss: 0.1923
Epoch [440/500], Loss: 0.1911
Epoch [450/500], Loss: 0.1890
Epoch [460/500], Loss: 0.1882
Epoch [470/500], Loss: 0.1963
Epoch [480/500], Loss: 0.1923
Epoch [490/500], Loss: 0.1878
Epoch [500/500], Loss: 0.1847
```

Train Accuracy: 0.9395

Test Accuracy: 0.9246

Try with lesser neurons

```
In [72]: first_layer_neurons = 512
         dropout_rate = 0.4
```

```

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.fc1 = nn.Linear(input_size, first_layer_neurons)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(first_layer_neurons, first_layer_neurons // 2)
        self.relu2 = nn.ReLU()
        self.fc3 = nn.Linear(first_layer_neurons // 2, first_layer_neurons // 4)
        self.relu3 = nn.ReLU()
        self.fc4 = nn.Linear(first_layer_neurons // 4, first_layer_neurons // 8)
        self.relu4 = nn.ReLU()
        self.fc5 = nn.Linear(first_layer_neurons // 8, 1)
        self.relu5 = nn.ReLU()
        self.dropout1 = nn.Dropout(dropout_rate)
        self.dropout2 = nn.Dropout(dropout_rate)
        self.dropout3 = nn.Dropout(dropout_rate)
        self.dropout4 = nn.Dropout(dropout_rate)
        self.dropout5 = nn.Dropout(dropout_rate)

    def forward(self, x):
        x = self.relu1(self.fc1(x))
        x = self.dropout1(x)
        x = self.relu2(self.fc2(x))
        x = self.dropout2(x)
        x = self.relu3(self.fc3(x))
        x = self.dropout3(x)
        x = self.relu4(self.fc4(x))
        x = self.dropout4(x)
        x = self.fc5(x)
        return x

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 500
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_test_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_test_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()

```

```
accuracy = (predicted_classes == y_test_tensor).float().mean()
print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

```
Epoch [10/500], Loss: 0.5262
Epoch [20/500], Loss: 0.4100
Epoch [30/500], Loss: 0.3815
Epoch [40/500], Loss: 0.3523
Epoch [50/500], Loss: 0.3227
Epoch [60/500], Loss: 0.3045
Epoch [70/500], Loss: 0.2848
Epoch [80/500], Loss: 0.2775
Epoch [90/500], Loss: 0.2700
Epoch [100/500], Loss: 0.2631
Epoch [110/500], Loss: 0.2565
Epoch [120/500], Loss: 0.2525
Epoch [130/500], Loss: 0.2458
Epoch [140/500], Loss: 0.2427
Epoch [150/500], Loss: 0.2393
Epoch [160/500], Loss: 0.2347
Epoch [170/500], Loss: 0.2323
Epoch [180/500], Loss: 0.2308
Epoch [190/500], Loss: 0.2294
Epoch [200/500], Loss: 0.2266
Epoch [210/500], Loss: 0.2249
Epoch [220/500], Loss: 0.2242
Epoch [230/500], Loss: 0.2217
Epoch [240/500], Loss: 0.2196
Epoch [250/500], Loss: 0.2201
Epoch [260/500], Loss: 0.2175
Epoch [270/500], Loss: 0.2138
Epoch [280/500], Loss: 0.2157
Epoch [290/500], Loss: 0.2134
Epoch [300/500], Loss: 0.2145
Epoch [310/500], Loss: 0.2148
Epoch [320/500], Loss: 0.2143
Epoch [330/500], Loss: 0.2109
Epoch [340/500], Loss: 0.2103
Epoch [350/500], Loss: 0.2117
Epoch [360/500], Loss: 0.2079
Epoch [370/500], Loss: 0.2097
Epoch [380/500], Loss: 0.2073
Epoch [390/500], Loss: 0.2081
Epoch [400/500], Loss: 0.2056
Epoch [410/500], Loss: 0.2047
Epoch [420/500], Loss: 0.2043
Epoch [430/500], Loss: 0.2050
Epoch [440/500], Loss: 0.2034
Epoch [450/500], Loss: 0.2029
Epoch [460/500], Loss: 0.2014
Epoch [470/500], Loss: 0.2021
Epoch [480/500], Loss: 0.2030
Epoch [490/500], Loss: 0.2035
Epoch [500/500], Loss: 0.2019
```

Train Accuracy: 0.9367

Test Accuracy: 0.9261

```
In [73]: first_layer_neurons = 256
         dropout_rate = 0.3
```

```

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            # nn.Linear(, 256),
            # nn.BatchNorm1d(256),
            # nn.ReLU(),
            # nn.Dropout(dropout_rate),
            nn.Linear(input_size, 128),
            nn.BatchNorm1d(128),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(128, 64),
            nn.BatchNorm1d(64),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(64, 32),
            nn.BatchNorm1d(32),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(32, 1),
        )

    def forward(self, x):
        return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 500
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")

```

```
Epoch [10/500], Loss: 0.7350
Epoch [20/500], Loss: 0.6287
Epoch [30/500], Loss: 0.5676
Epoch [40/500], Loss: 0.5266
Epoch [50/500], Loss: 0.4984
Epoch [60/500], Loss: 0.4739
Epoch [70/500], Loss: 0.4519
Epoch [80/500], Loss: 0.4338
Epoch [90/500], Loss: 0.4138
Epoch [100/500], Loss: 0.3993
Epoch [110/500], Loss: 0.3836
Epoch [120/500], Loss: 0.3690
Epoch [130/500], Loss: 0.3554
Epoch [140/500], Loss: 0.3449
Epoch [150/500], Loss: 0.3342
Epoch [160/500], Loss: 0.3255
Epoch [170/500], Loss: 0.3162
Epoch [180/500], Loss: 0.3052
Epoch [190/500], Loss: 0.3022
Epoch [200/500], Loss: 0.2958
Epoch [210/500], Loss: 0.2909
Epoch [220/500], Loss: 0.2866
Epoch [230/500], Loss: 0.2790
Epoch [240/500], Loss: 0.2767
Epoch [250/500], Loss: 0.2702
Epoch [260/500], Loss: 0.2702
Epoch [270/500], Loss: 0.2637
Epoch [280/500], Loss: 0.2627
Epoch [290/500], Loss: 0.2596
Epoch [300/500], Loss: 0.2561
Epoch [310/500], Loss: 0.2528
Epoch [320/500], Loss: 0.2501
Epoch [330/500], Loss: 0.2482
Epoch [340/500], Loss: 0.2460
Epoch [350/500], Loss: 0.2428
Epoch [360/500], Loss: 0.2415
Epoch [370/500], Loss: 0.2397
Epoch [380/500], Loss: 0.2391
Epoch [390/500], Loss: 0.2388
Epoch [400/500], Loss: 0.2340
Epoch [410/500], Loss: 0.2310
Epoch [420/500], Loss: 0.2298
Epoch [430/500], Loss: 0.2269
Epoch [440/500], Loss: 0.2289
Epoch [450/500], Loss: 0.2220
Epoch [460/500], Loss: 0.2220
Epoch [470/500], Loss: 0.2248
Epoch [480/500], Loss: 0.2199
Epoch [490/500], Loss: 0.2212
Epoch [500/500], Loss: 0.2184
```

Train Accuracy: 0.9302

Test Accuracy: 0.9249

```
In [74]: first_layer_neurons = 256
         dropout_rate = 0.3

class LoanDefaultNN(nn.Module):
```



```

def __init__(self, input_size):
    super(LoanDefaultNN, self).__init__()
    self.model = nn.Sequential(
        # nn.Linear(, 256),
        # nn.BatchNorm1d(256),
        # nn.ReLU(),
        # nn.Dropout(dropout_rate),
        nn.Linear(input_size, 26),
        nn.BatchNorm1d(26),
        nn.LeakyReLU(0.01),
        nn.Dropout(dropout_rate),
        nn.Linear(26, 13),
        nn.BatchNorm1d(13),
        nn.LeakyReLU(0.01),
        nn.Dropout(dropout_rate),
        nn.Linear(13, 6),
        nn.BatchNorm1d(6),
        nn.LeakyReLU(0.01),
        nn.Dropout(dropout_rate),
        nn.Linear(6, 1),
    )

def forward(self, x):
    return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 2000
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")

```

Epoch [10/2000], Loss: 0.6066  
Epoch [20/2000], Loss: 0.5785  
Epoch [30/2000], Loss: 0.5595  
Epoch [40/2000], Loss: 0.5420  
Epoch [50/2000], Loss: 0.5271  
Epoch [60/2000], Loss: 0.5143  
Epoch [70/2000], Loss: 0.5022  
Epoch [80/2000], Loss: 0.4888  
Epoch [90/2000], Loss: 0.4791  
Epoch [100/2000], Loss: 0.4721  
Epoch [110/2000], Loss: 0.4596  
Epoch [120/2000], Loss: 0.4519  
Epoch [130/2000], Loss: 0.4449  
Epoch [140/2000], Loss: 0.4365  
Epoch [150/2000], Loss: 0.4313  
Epoch [160/2000], Loss: 0.4245  
Epoch [170/2000], Loss: 0.4148  
Epoch [180/2000], Loss: 0.4115  
Epoch [190/2000], Loss: 0.4052  
Epoch [200/2000], Loss: 0.3971  
Epoch [210/2000], Loss: 0.3942  
Epoch [220/2000], Loss: 0.3890  
Epoch [230/2000], Loss: 0.3830  
Epoch [240/2000], Loss: 0.3777  
Epoch [250/2000], Loss: 0.3749  
Epoch [260/2000], Loss: 0.3710  
Epoch [270/2000], Loss: 0.3664  
Epoch [280/2000], Loss: 0.3589  
Epoch [290/2000], Loss: 0.3588  
Epoch [300/2000], Loss: 0.3546  
Epoch [310/2000], Loss: 0.3531  
Epoch [320/2000], Loss: 0.3467  
Epoch [330/2000], Loss: 0.3447  
Epoch [340/2000], Loss: 0.3390  
Epoch [350/2000], Loss: 0.3363  
Epoch [360/2000], Loss: 0.3331  
Epoch [370/2000], Loss: 0.3313  
Epoch [380/2000], Loss: 0.3253  
Epoch [390/2000], Loss: 0.3212  
Epoch [400/2000], Loss: 0.3170  
Epoch [410/2000], Loss: 0.3170  
Epoch [420/2000], Loss: 0.3140  
Epoch [430/2000], Loss: 0.3087  
Epoch [440/2000], Loss: 0.3070  
Epoch [450/2000], Loss: 0.3033  
Epoch [460/2000], Loss: 0.3031  
Epoch [470/2000], Loss: 0.2977  
Epoch [480/2000], Loss: 0.2962  
Epoch [490/2000], Loss: 0.2953  
Epoch [500/2000], Loss: 0.2945  
Epoch [510/2000], Loss: 0.2891  
Epoch [520/2000], Loss: 0.2889  
Epoch [530/2000], Loss: 0.2847  
Epoch [540/2000], Loss: 0.2837  
Epoch [550/2000], Loss: 0.2836  
Epoch [560/2000], Loss: 0.2837  
Epoch [570/2000], Loss: 0.2789  
Epoch [580/2000], Loss: 0.2824  
Epoch [590/2000], Loss: 0.2742  
Epoch [600/2000], Loss: 0.2769

Epoch [610/2000], Loss: 0.2801  
Epoch [620/2000], Loss: 0.2739  
Epoch [630/2000], Loss: 0.2742  
Epoch [640/2000], Loss: 0.2753  
Epoch [650/2000], Loss: 0.2740  
Epoch [660/2000], Loss: 0.2711  
Epoch [670/2000], Loss: 0.2715  
Epoch [680/2000], Loss: 0.2706  
Epoch [690/2000], Loss: 0.2684  
Epoch [700/2000], Loss: 0.2689  
Epoch [710/2000], Loss: 0.2707  
Epoch [720/2000], Loss: 0.2692  
Epoch [730/2000], Loss: 0.2686  
Epoch [740/2000], Loss: 0.2688  
Epoch [750/2000], Loss: 0.2680  
Epoch [760/2000], Loss: 0.2686  
Epoch [770/2000], Loss: 0.2664  
Epoch [780/2000], Loss: 0.2685  
Epoch [790/2000], Loss: 0.2644  
Epoch [800/2000], Loss: 0.2638  
Epoch [810/2000], Loss: 0.2644  
Epoch [820/2000], Loss: 0.2625  
Epoch [830/2000], Loss: 0.2617  
Epoch [840/2000], Loss: 0.2598  
Epoch [850/2000], Loss: 0.2593  
Epoch [860/2000], Loss: 0.2646  
Epoch [870/2000], Loss: 0.2633  
Epoch [880/2000], Loss: 0.2581  
Epoch [890/2000], Loss: 0.2615  
Epoch [900/2000], Loss: 0.2614  
Epoch [910/2000], Loss: 0.2601  
Epoch [920/2000], Loss: 0.2617  
Epoch [930/2000], Loss: 0.2591  
Epoch [940/2000], Loss: 0.2594  
Epoch [950/2000], Loss: 0.2602  
Epoch [960/2000], Loss: 0.2563  
Epoch [970/2000], Loss: 0.2599  
Epoch [980/2000], Loss: 0.2583  
Epoch [990/2000], Loss: 0.2570  
Epoch [1000/2000], Loss: 0.2568  
Epoch [1010/2000], Loss: 0.2553  
Epoch [1020/2000], Loss: 0.2557  
Epoch [1030/2000], Loss: 0.2557  
Epoch [1040/2000], Loss: 0.2534  
Epoch [1050/2000], Loss: 0.2592  
Epoch [1060/2000], Loss: 0.2568  
Epoch [1070/2000], Loss: 0.2538  
Epoch [1080/2000], Loss: 0.2560  
Epoch [1090/2000], Loss: 0.2569  
Epoch [1100/2000], Loss: 0.2517  
Epoch [1110/2000], Loss: 0.2511  
Epoch [1120/2000], Loss: 0.2510  
Epoch [1130/2000], Loss: 0.2545  
Epoch [1140/2000], Loss: 0.2526  
Epoch [1150/2000], Loss: 0.2528  
Epoch [1160/2000], Loss: 0.2545  
Epoch [1170/2000], Loss: 0.2532  
Epoch [1180/2000], Loss: 0.2519  
Epoch [1190/2000], Loss: 0.2541  
Epoch [1200/2000], Loss: 0.2535

Epoch [1210/2000], Loss: 0.2525  
Epoch [1220/2000], Loss: 0.2507  
Epoch [1230/2000], Loss: 0.2508  
Epoch [1240/2000], Loss: 0.2517  
Epoch [1250/2000], Loss: 0.2524  
Epoch [1260/2000], Loss: 0.2503  
Epoch [1270/2000], Loss: 0.2512  
Epoch [1280/2000], Loss: 0.2497  
Epoch [1290/2000], Loss: 0.2515  
Epoch [1300/2000], Loss: 0.2544  
Epoch [1310/2000], Loss: 0.2515  
Epoch [1320/2000], Loss: 0.2510  
Epoch [1330/2000], Loss: 0.2492  
Epoch [1340/2000], Loss: 0.2530  
Epoch [1350/2000], Loss: 0.2498  
Epoch [1360/2000], Loss: 0.2478  
Epoch [1370/2000], Loss: 0.2502  
Epoch [1380/2000], Loss: 0.2485  
Epoch [1390/2000], Loss: 0.2507  
Epoch [1400/2000], Loss: 0.2495  
Epoch [1410/2000], Loss: 0.2484  
Epoch [1420/2000], Loss: 0.2469  
Epoch [1430/2000], Loss: 0.2486  
Epoch [1440/2000], Loss: 0.2477  
Epoch [1450/2000], Loss: 0.2471  
Epoch [1460/2000], Loss: 0.2467  
Epoch [1470/2000], Loss: 0.2466  
Epoch [1480/2000], Loss: 0.2481  
Epoch [1490/2000], Loss: 0.2462  
Epoch [1500/2000], Loss: 0.2484  
Epoch [1510/2000], Loss: 0.2480  
Epoch [1520/2000], Loss: 0.2472  
Epoch [1530/2000], Loss: 0.2481  
Epoch [1540/2000], Loss: 0.2510  
Epoch [1550/2000], Loss: 0.2464  
Epoch [1560/2000], Loss: 0.2479  
Epoch [1570/2000], Loss: 0.2476  
Epoch [1580/2000], Loss: 0.2477  
Epoch [1590/2000], Loss: 0.2463  
Epoch [1600/2000], Loss: 0.2489  
Epoch [1610/2000], Loss: 0.2464  
Epoch [1620/2000], Loss: 0.2496  
Epoch [1630/2000], Loss: 0.2470  
Epoch [1640/2000], Loss: 0.2508  
Epoch [1650/2000], Loss: 0.2475  
Epoch [1660/2000], Loss: 0.2500  
Epoch [1670/2000], Loss: 0.2475  
Epoch [1680/2000], Loss: 0.2449  
Epoch [1690/2000], Loss: 0.2472  
Epoch [1700/2000], Loss: 0.2500  
Epoch [1710/2000], Loss: 0.2469  
Epoch [1720/2000], Loss: 0.2453  
Epoch [1730/2000], Loss: 0.2453  
Epoch [1740/2000], Loss: 0.2483  
Epoch [1750/2000], Loss: 0.2449  
Epoch [1760/2000], Loss: 0.2433  
Epoch [1770/2000], Loss: 0.2457  
Epoch [1780/2000], Loss: 0.2450  
Epoch [1790/2000], Loss: 0.2464  
Epoch [1800/2000], Loss: 0.2475

```
Epoch [1810/2000], Loss: 0.2477
Epoch [1820/2000], Loss: 0.2430
Epoch [1830/2000], Loss: 0.2461
Epoch [1840/2000], Loss: 0.2471
Epoch [1850/2000], Loss: 0.2469
Epoch [1860/2000], Loss: 0.2453
Epoch [1870/2000], Loss: 0.2468
Epoch [1880/2000], Loss: 0.2471
Epoch [1890/2000], Loss: 0.2454
Epoch [1900/2000], Loss: 0.2481
Epoch [1910/2000], Loss: 0.2457
Epoch [1920/2000], Loss: 0.2441
Epoch [1930/2000], Loss: 0.2474
Epoch [1940/2000], Loss: 0.2460
Epoch [1950/2000], Loss: 0.2472
Epoch [1960/2000], Loss: 0.2467
Epoch [1970/2000], Loss: 0.2444
Epoch [1980/2000], Loss: 0.2407
Epoch [1990/2000], Loss: 0.2449
Epoch [2000/2000], Loss: 0.2450
```

Train Accuracy: 0.9321

Test Accuracy: 0.9282

```
In [75]: first_layer_neurons = 256
         dropout_rate = 0.2

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            # nn.Linear(, 256),
            # nn.BatchNorm1d(256),
            # nn.ReLU(),
            # nn.Dropout(dropout_rate),
            nn.Linear(input_size, 26),
            nn.BatchNorm1d(26),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(26, 13),
            nn.BatchNorm1d(13),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(13, 6),
            nn.BatchNorm1d(6),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(6, 1),
        )

    def forward(self, x):
        return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 2000
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/2000], Loss: 0.8835  
Epoch [20/2000], Loss: 0.8384  
Epoch [30/2000], Loss: 0.8035  
Epoch [40/2000], Loss: 0.7787  
Epoch [50/2000], Loss: 0.7556  
Epoch [60/2000], Loss: 0.7368  
Epoch [70/2000], Loss: 0.7202  
Epoch [80/2000], Loss: 0.7051  
Epoch [90/2000], Loss: 0.6904  
Epoch [100/2000], Loss: 0.6752  
Epoch [110/2000], Loss: 0.6604  
Epoch [120/2000], Loss: 0.6456  
Epoch [130/2000], Loss: 0.6327  
Epoch [140/2000], Loss: 0.6172  
Epoch [150/2000], Loss: 0.6016  
Epoch [160/2000], Loss: 0.5894  
Epoch [170/2000], Loss: 0.5728  
Epoch [180/2000], Loss: 0.5602  
Epoch [190/2000], Loss: 0.5470  
Epoch [200/2000], Loss: 0.5357  
Epoch [210/2000], Loss: 0.5211  
Epoch [220/2000], Loss: 0.5110  
Epoch [230/2000], Loss: 0.4979  
Epoch [240/2000], Loss: 0.4868  
Epoch [250/2000], Loss: 0.4750  
Epoch [260/2000], Loss: 0.4654  
Epoch [270/2000], Loss: 0.4562  
Epoch [280/2000], Loss: 0.4488  
Epoch [290/2000], Loss: 0.4392  
Epoch [300/2000], Loss: 0.4341  
Epoch [310/2000], Loss: 0.4229  
Epoch [320/2000], Loss: 0.4177  
Epoch [330/2000], Loss: 0.4131  
Epoch [340/2000], Loss: 0.4042  
Epoch [350/2000], Loss: 0.3980  
Epoch [360/2000], Loss: 0.3925  
Epoch [370/2000], Loss: 0.3891  
Epoch [380/2000], Loss: 0.3846  
Epoch [390/2000], Loss: 0.3791  
Epoch [400/2000], Loss: 0.3743  
Epoch [410/2000], Loss: 0.3707  
Epoch [420/2000], Loss: 0.3672  
Epoch [430/2000], Loss: 0.3636  
Epoch [440/2000], Loss: 0.3596  
Epoch [450/2000], Loss: 0.3537  
Epoch [460/2000], Loss: 0.3507  
Epoch [470/2000], Loss: 0.3489  
Epoch [480/2000], Loss: 0.3431  
Epoch [490/2000], Loss: 0.3389  
Epoch [500/2000], Loss: 0.3292  
Epoch [510/2000], Loss: 0.3196  
Epoch [520/2000], Loss: 0.3169  
Epoch [530/2000], Loss: 0.3115  
Epoch [540/2000], Loss: 0.3055  
Epoch [550/2000], Loss: 0.3063  
Epoch [560/2000], Loss: 0.2981  
Epoch [570/2000], Loss: 0.2962  
Epoch [580/2000], Loss: 0.2933  
Epoch [590/2000], Loss: 0.2925  
Epoch [600/2000], Loss: 0.2866

Epoch [610/2000], Loss: 0.2868  
Epoch [620/2000], Loss: 0.2816  
Epoch [630/2000], Loss: 0.2834  
Epoch [640/2000], Loss: 0.2801  
Epoch [650/2000], Loss: 0.2777  
Epoch [660/2000], Loss: 0.2769  
Epoch [670/2000], Loss: 0.2743  
Epoch [680/2000], Loss: 0.2728  
Epoch [690/2000], Loss: 0.2713  
Epoch [700/2000], Loss: 0.2700  
Epoch [710/2000], Loss: 0.2720  
Epoch [720/2000], Loss: 0.2674  
Epoch [730/2000], Loss: 0.2650  
Epoch [740/2000], Loss: 0.2643  
Epoch [750/2000], Loss: 0.2663  
Epoch [760/2000], Loss: 0.2609  
Epoch [770/2000], Loss: 0.2613  
Epoch [780/2000], Loss: 0.2578  
Epoch [790/2000], Loss: 0.2592  
Epoch [800/2000], Loss: 0.2561  
Epoch [810/2000], Loss: 0.2538  
Epoch [820/2000], Loss: 0.2538  
Epoch [830/2000], Loss: 0.2543  
Epoch [840/2000], Loss: 0.2512  
Epoch [850/2000], Loss: 0.2533  
Epoch [860/2000], Loss: 0.2503  
Epoch [870/2000], Loss: 0.2522  
Epoch [880/2000], Loss: 0.2507  
Epoch [890/2000], Loss: 0.2480  
Epoch [900/2000], Loss: 0.2470  
Epoch [910/2000], Loss: 0.2476  
Epoch [920/2000], Loss: 0.2466  
Epoch [930/2000], Loss: 0.2460  
Epoch [940/2000], Loss: 0.2440  
Epoch [950/2000], Loss: 0.2453  
Epoch [960/2000], Loss: 0.2442  
Epoch [970/2000], Loss: 0.2430  
Epoch [980/2000], Loss: 0.2434  
Epoch [990/2000], Loss: 0.2447  
Epoch [1000/2000], Loss: 0.2403  
Epoch [1010/2000], Loss: 0.2408  
Epoch [1020/2000], Loss: 0.2405  
Epoch [1030/2000], Loss: 0.2428  
Epoch [1040/2000], Loss: 0.2404  
Epoch [1050/2000], Loss: 0.2431  
Epoch [1060/2000], Loss: 0.2413  
Epoch [1070/2000], Loss: 0.2395  
Epoch [1080/2000], Loss: 0.2397  
Epoch [1090/2000], Loss: 0.2395  
Epoch [1100/2000], Loss: 0.2381  
Epoch [1110/2000], Loss: 0.2408  
Epoch [1120/2000], Loss: 0.2399  
Epoch [1130/2000], Loss: 0.2405  
Epoch [1140/2000], Loss: 0.2391  
Epoch [1150/2000], Loss: 0.2397  
Epoch [1160/2000], Loss: 0.2386  
Epoch [1170/2000], Loss: 0.2360  
Epoch [1180/2000], Loss: 0.2404  
Epoch [1190/2000], Loss: 0.2380  
Epoch [1200/2000], Loss: 0.2354



Epoch [1210/2000], Loss: 0.2361  
Epoch [1220/2000], Loss: 0.2389  
Epoch [1230/2000], Loss: 0.2356  
Epoch [1240/2000], Loss: 0.2361  
Epoch [1250/2000], Loss: 0.2357  
Epoch [1260/2000], Loss: 0.2359  
Epoch [1270/2000], Loss: 0.2338  
Epoch [1280/2000], Loss: 0.2401  
Epoch [1290/2000], Loss: 0.2326  
Epoch [1300/2000], Loss: 0.2369  
Epoch [1310/2000], Loss: 0.2323  
Epoch [1320/2000], Loss: 0.2345  
Epoch [1330/2000], Loss: 0.2324  
Epoch [1340/2000], Loss: 0.2346  
Epoch [1350/2000], Loss: 0.2336  
Epoch [1360/2000], Loss: 0.2354  
Epoch [1370/2000], Loss: 0.2350  
Epoch [1380/2000], Loss: 0.2329  
Epoch [1390/2000], Loss: 0.2335  
Epoch [1400/2000], Loss: 0.2332  
Epoch [1410/2000], Loss: 0.2346  
Epoch [1420/2000], Loss: 0.2352  
Epoch [1430/2000], Loss: 0.2367  
Epoch [1440/2000], Loss: 0.2342  
Epoch [1450/2000], Loss: 0.2352  
Epoch [1460/2000], Loss: 0.2329  
Epoch [1470/2000], Loss: 0.2340  
Epoch [1480/2000], Loss: 0.2323  
Epoch [1490/2000], Loss: 0.2314  
Epoch [1500/2000], Loss: 0.2334  
Epoch [1510/2000], Loss: 0.2325  
Epoch [1520/2000], Loss: 0.2303  
Epoch [1530/2000], Loss: 0.2321  
Epoch [1540/2000], Loss: 0.2331  
Epoch [1550/2000], Loss: 0.2317  
Epoch [1560/2000], Loss: 0.2336  
Epoch [1570/2000], Loss: 0.2331  
Epoch [1580/2000], Loss: 0.2328  
Epoch [1590/2000], Loss: 0.2300  
Epoch [1600/2000], Loss: 0.2327  
Epoch [1610/2000], Loss: 0.2345  
Epoch [1620/2000], Loss: 0.2314  
Epoch [1630/2000], Loss: 0.2377  
Epoch [1640/2000], Loss: 0.2335  
Epoch [1650/2000], Loss: 0.2318  
Epoch [1660/2000], Loss: 0.2312  
Epoch [1670/2000], Loss: 0.2327  
Epoch [1680/2000], Loss: 0.2292  
Epoch [1690/2000], Loss: 0.2320  
Epoch [1700/2000], Loss: 0.2323  
Epoch [1710/2000], Loss: 0.2330  
Epoch [1720/2000], Loss: 0.2314  
Epoch [1730/2000], Loss: 0.2326  
Epoch [1740/2000], Loss: 0.2327  
Epoch [1750/2000], Loss: 0.2291  
Epoch [1760/2000], Loss: 0.2317  
Epoch [1770/2000], Loss: 0.2303  
Epoch [1780/2000], Loss: 0.2344  
Epoch [1790/2000], Loss: 0.2325  
Epoch [1800/2000], Loss: 0.2292

```
Epoch [1810/2000], Loss: 0.2313
Epoch [1820/2000], Loss: 0.2324
Epoch [1830/2000], Loss: 0.2338
Epoch [1840/2000], Loss: 0.2330
Epoch [1850/2000], Loss: 0.2310
Epoch [1860/2000], Loss: 0.2293
Epoch [1870/2000], Loss: 0.2315
Epoch [1880/2000], Loss: 0.2290
Epoch [1890/2000], Loss: 0.2306
Epoch [1900/2000], Loss: 0.2281
Epoch [1910/2000], Loss: 0.2336
Epoch [1920/2000], Loss: 0.2283
Epoch [1930/2000], Loss: 0.2277
Epoch [1940/2000], Loss: 0.2308
Epoch [1950/2000], Loss: 0.2304
Epoch [1960/2000], Loss: 0.2316
Epoch [1970/2000], Loss: 0.2315
Epoch [1980/2000], Loss: 0.2286
Epoch [1990/2000], Loss: 0.2280
Epoch [2000/2000], Loss: 0.2290
```

Train Accuracy: 0.9339

Test Accuracy: 0.9293

```
In [76]: first_layer_neurons = 256
         dropout_rate = 0.3

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            # nn.Linear(, 256),
            # nn.BatchNorm1d(256),
            # nn.ReLU(),
            # nn.Dropout(dropout_rate),
            nn.Linear(input_size, 64),
            nn.BatchNorm1d(64),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(64, 32),
            nn.BatchNorm1d(32),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(32, 16),
            nn.BatchNorm1d(16),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(16, 8),
            nn.BatchNorm1d(8),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(8, 4),
            nn.BatchNorm1d(4),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(4, 1),
        )
```

```
def forward(self, x):  
    return self.model(x) # Use if you're sticking with BCELoss  
  
# Initialize model, loss, optimizer  
model = LoanDefaultNN(X_train.shape[1])  
criterion = nn.BCEWithLogitsLoss()  
# Add regularization in weight decay  
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)  
  
# Training loop  
epochs = 1000  
for epoch in range(epochs):  
    model.train()  
    optimizer.zero_grad()  
    output = model(X_train_tensor)  
    loss = criterion(output, y_train_tensor)  
    loss.backward()  
    optimizer.step()  
  
    if (epoch + 1) % 10 == 0:  
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")  
  
model.eval()  
with torch.no_grad():  
    test_preds = model(X_train_tensor)  
    test_predicted_classes = (test_preds > 0.5).float()  
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()  
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")  
    preds = model(X_test_tensor)  
    predicted_classes = (preds > 0.5).float()  
    accuracy = (predicted_classes == y_test_tensor).float().mean()  
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/1000], Loss: 0.7695  
Epoch [20/1000], Loss: 0.7368  
Epoch [30/1000], Loss: 0.7066  
Epoch [40/1000], Loss: 0.6839  
Epoch [50/1000], Loss: 0.6661  
Epoch [60/1000], Loss: 0.6514  
Epoch [70/1000], Loss: 0.6359  
Epoch [80/1000], Loss: 0.6255  
Epoch [90/1000], Loss: 0.6136  
Epoch [100/1000], Loss: 0.6021  
Epoch [110/1000], Loss: 0.5913  
Epoch [120/1000], Loss: 0.5828  
Epoch [130/1000], Loss: 0.5743  
Epoch [140/1000], Loss: 0.5632  
Epoch [150/1000], Loss: 0.5564  
Epoch [160/1000], Loss: 0.5501  
Epoch [170/1000], Loss: 0.5404  
Epoch [180/1000], Loss: 0.5336  
Epoch [190/1000], Loss: 0.5251  
Epoch [200/1000], Loss: 0.5194  
Epoch [210/1000], Loss: 0.5090  
Epoch [220/1000], Loss: 0.5043  
Epoch [230/1000], Loss: 0.4987  
Epoch [240/1000], Loss: 0.4946  
Epoch [250/1000], Loss: 0.4875  
Epoch [260/1000], Loss: 0.4788  
Epoch [270/1000], Loss: 0.4739  
Epoch [280/1000], Loss: 0.4675  
Epoch [290/1000], Loss: 0.4648  
Epoch [300/1000], Loss: 0.4586  
Epoch [310/1000], Loss: 0.4542  
Epoch [320/1000], Loss: 0.4490  
Epoch [330/1000], Loss: 0.4455  
Epoch [340/1000], Loss: 0.4386  
Epoch [350/1000], Loss: 0.4359  
Epoch [360/1000], Loss: 0.4250  
Epoch [370/1000], Loss: 0.4233  
Epoch [380/1000], Loss: 0.4218  
Epoch [390/1000], Loss: 0.4167  
Epoch [400/1000], Loss: 0.4085  
Epoch [410/1000], Loss: 0.4055  
Epoch [420/1000], Loss: 0.4010  
Epoch [430/1000], Loss: 0.3994  
Epoch [440/1000], Loss: 0.3960  
Epoch [450/1000], Loss: 0.3917  
Epoch [460/1000], Loss: 0.3884  
Epoch [470/1000], Loss: 0.3853  
Epoch [480/1000], Loss: 0.3832  
Epoch [490/1000], Loss: 0.3775  
Epoch [500/1000], Loss: 0.3734  
Epoch [510/1000], Loss: 0.3715  
Epoch [520/1000], Loss: 0.3670  
Epoch [530/1000], Loss: 0.3667  
Epoch [540/1000], Loss: 0.3634  
Epoch [550/1000], Loss: 0.3619  
Epoch [560/1000], Loss: 0.3622  
Epoch [570/1000], Loss: 0.3557  
Epoch [580/1000], Loss: 0.3568  
Epoch [590/1000], Loss: 0.3529  
Epoch [600/1000], Loss: 0.3514

```
Epoch [610/1000], Loss: 0.3499
Epoch [620/1000], Loss: 0.3413
Epoch [630/1000], Loss: 0.3400
Epoch [640/1000], Loss: 0.3402
Epoch [650/1000], Loss: 0.3413
Epoch [660/1000], Loss: 0.3375
Epoch [670/1000], Loss: 0.3358
Epoch [680/1000], Loss: 0.3334
Epoch [690/1000], Loss: 0.3328
Epoch [700/1000], Loss: 0.3321
Epoch [710/1000], Loss: 0.3275
Epoch [720/1000], Loss: 0.3270
Epoch [730/1000], Loss: 0.3234
Epoch [740/1000], Loss: 0.3209
Epoch [750/1000], Loss: 0.3227
Epoch [760/1000], Loss: 0.3206
Epoch [770/1000], Loss: 0.3223
Epoch [780/1000], Loss: 0.3182
Epoch [790/1000], Loss: 0.3205
Epoch [800/1000], Loss: 0.3181
Epoch [810/1000], Loss: 0.3144
Epoch [820/1000], Loss: 0.3155
Epoch [830/1000], Loss: 0.3147
Epoch [840/1000], Loss: 0.3135
Epoch [850/1000], Loss: 0.3135
Epoch [860/1000], Loss: 0.3114
Epoch [870/1000], Loss: 0.3122
Epoch [880/1000], Loss: 0.3105
Epoch [890/1000], Loss: 0.3122
Epoch [900/1000], Loss: 0.3099
Epoch [910/1000], Loss: 0.3042
Epoch [920/1000], Loss: 0.3082
Epoch [930/1000], Loss: 0.3050
Epoch [940/1000], Loss: 0.3087
Epoch [950/1000], Loss: 0.3027
Epoch [960/1000], Loss: 0.3015
Epoch [970/1000], Loss: 0.3040
Epoch [980/1000], Loss: 0.3027
Epoch [990/1000], Loss: 0.3019
Epoch [1000/1000], Loss: 0.3008
```

Train Accuracy: 0.9364

Test Accuracy: 0.9302

```
In [77]: first_layer_neurons = 256
         dropout_rate = 0.3

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(input_size, 52),
            nn.BatchNorm1d(52),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(52, 26),
            nn.BatchNorm1d(26),
            nn.LeakyReLU(0.01),
```

```

        nn.Dropout(dropout_rate),
        nn.Linear(26, 13),
        nn.BatchNorm1d(13),
        nn.LeakyReLU(0.01),
        nn.Dropout(dropout_rate),
        nn.Linear(13, 6),
        nn.BatchNorm1d(6),
        nn.LeakyReLU(0.01),
        nn.Dropout(dropout_rate),
        # nn.Linear(6, 3),
        # nn.BatchNorm1d(3),
        # nn.LeakyReLU(0.01),
        # nn.Dropout(dropout_rate),
        nn.Linear(6, 1),
    )

    def forward(self, x):
        return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 1500
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")

```

Epoch [10/1500], Loss: 0.6404  
Epoch [20/1500], Loss: 0.6045  
Epoch [30/1500], Loss: 0.5781  
Epoch [40/1500], Loss: 0.5572  
Epoch [50/1500], Loss: 0.5399  
Epoch [60/1500], Loss: 0.5214  
Epoch [70/1500], Loss: 0.5088  
Epoch [80/1500], Loss: 0.4972  
Epoch [90/1500], Loss: 0.4843  
Epoch [100/1500], Loss: 0.4705  
Epoch [110/1500], Loss: 0.4596  
Epoch [120/1500], Loss: 0.4509  
Epoch [130/1500], Loss: 0.4404  
Epoch [140/1500], Loss: 0.4324  
Epoch [150/1500], Loss: 0.4260  
Epoch [160/1500], Loss: 0.4140  
Epoch [170/1500], Loss: 0.4095  
Epoch [180/1500], Loss: 0.3994  
Epoch [190/1500], Loss: 0.3937  
Epoch [200/1500], Loss: 0.3859  
Epoch [210/1500], Loss: 0.3788  
Epoch [220/1500], Loss: 0.3731  
Epoch [230/1500], Loss: 0.3710  
Epoch [240/1500], Loss: 0.3601  
Epoch [250/1500], Loss: 0.3601  
Epoch [260/1500], Loss: 0.3558  
Epoch [270/1500], Loss: 0.3479  
Epoch [280/1500], Loss: 0.3445  
Epoch [290/1500], Loss: 0.3415  
Epoch [300/1500], Loss: 0.3382  
Epoch [310/1500], Loss: 0.3365  
Epoch [320/1500], Loss: 0.3314  
Epoch [330/1500], Loss: 0.3281  
Epoch [340/1500], Loss: 0.3267  
Epoch [350/1500], Loss: 0.3228  
Epoch [360/1500], Loss: 0.3193  
Epoch [370/1500], Loss: 0.3184  
Epoch [380/1500], Loss: 0.3143  
Epoch [390/1500], Loss: 0.3112  
Epoch [400/1500], Loss: 0.3119  
Epoch [410/1500], Loss: 0.3105  
Epoch [420/1500], Loss: 0.3060  
Epoch [430/1500], Loss: 0.3028  
Epoch [440/1500], Loss: 0.3057  
Epoch [450/1500], Loss: 0.3024  
Epoch [460/1500], Loss: 0.2968  
Epoch [470/1500], Loss: 0.2988  
Epoch [480/1500], Loss: 0.2947  
Epoch [490/1500], Loss: 0.2920  
Epoch [500/1500], Loss: 0.2898  
Epoch [510/1500], Loss: 0.2884  
Epoch [520/1500], Loss: 0.2877  
Epoch [530/1500], Loss: 0.2835  
Epoch [540/1500], Loss: 0.2836  
Epoch [550/1500], Loss: 0.2824  
Epoch [560/1500], Loss: 0.2841  
Epoch [570/1500], Loss: 0.2812  
Epoch [580/1500], Loss: 0.2775  
Epoch [590/1500], Loss: 0.2768  
Epoch [600/1500], Loss: 0.2789

Epoch [610/1500], Loss: 0.2782  
Epoch [620/1500], Loss: 0.2758  
Epoch [630/1500], Loss: 0.2756  
Epoch [640/1500], Loss: 0.2759  
Epoch [650/1500], Loss: 0.2723  
Epoch [660/1500], Loss: 0.2728  
Epoch [670/1500], Loss: 0.2718  
Epoch [680/1500], Loss: 0.2714  
Epoch [690/1500], Loss: 0.2719  
Epoch [700/1500], Loss: 0.2691  
Epoch [710/1500], Loss: 0.2715  
Epoch [720/1500], Loss: 0.2677  
Epoch [730/1500], Loss: 0.2705  
Epoch [740/1500], Loss: 0.2710  
Epoch [750/1500], Loss: 0.2656  
Epoch [760/1500], Loss: 0.2696  
Epoch [770/1500], Loss: 0.2665  
Epoch [780/1500], Loss: 0.2646  
Epoch [790/1500], Loss: 0.2648  
Epoch [800/1500], Loss: 0.2646  
Epoch [810/1500], Loss: 0.2662  
Epoch [820/1500], Loss: 0.2646  
Epoch [830/1500], Loss: 0.2671  
Epoch [840/1500], Loss: 0.2623  
Epoch [850/1500], Loss: 0.2639  
Epoch [860/1500], Loss: 0.2634  
Epoch [870/1500], Loss: 0.2626  
Epoch [880/1500], Loss: 0.2630  
Epoch [890/1500], Loss: 0.2610  
Epoch [900/1500], Loss: 0.2614  
Epoch [910/1500], Loss: 0.2646  
Epoch [920/1500], Loss: 0.2603  
Epoch [930/1500], Loss: 0.2617  
Epoch [940/1500], Loss: 0.2599  
Epoch [950/1500], Loss: 0.2603  
Epoch [960/1500], Loss: 0.2584  
Epoch [970/1500], Loss: 0.2584  
Epoch [980/1500], Loss: 0.2589  
Epoch [990/1500], Loss: 0.2593  
Epoch [1000/1500], Loss: 0.2614  
Epoch [1010/1500], Loss: 0.2592  
Epoch [1020/1500], Loss: 0.2592  
Epoch [1030/1500], Loss: 0.2588  
Epoch [1040/1500], Loss: 0.2583  
Epoch [1050/1500], Loss: 0.2587  
Epoch [1060/1500], Loss: 0.2595  
Epoch [1070/1500], Loss: 0.2575  
Epoch [1080/1500], Loss: 0.2552  
Epoch [1090/1500], Loss: 0.2584  
Epoch [1100/1500], Loss: 0.2585  
Epoch [1110/1500], Loss: 0.2541  
Epoch [1120/1500], Loss: 0.2581  
Epoch [1130/1500], Loss: 0.2580  
Epoch [1140/1500], Loss: 0.2561  
Epoch [1150/1500], Loss: 0.2565  
Epoch [1160/1500], Loss: 0.2549  
Epoch [1170/1500], Loss: 0.2556  
Epoch [1180/1500], Loss: 0.2542  
Epoch [1190/1500], Loss: 0.2550  
Epoch [1200/1500], Loss: 0.2565



```
Epoch [1210/1500], Loss: 0.2530
Epoch [1220/1500], Loss: 0.2556
Epoch [1230/1500], Loss: 0.2523
Epoch [1240/1500], Loss: 0.2561
Epoch [1250/1500], Loss: 0.2553
Epoch [1260/1500], Loss: 0.2541
Epoch [1270/1500], Loss: 0.2561
Epoch [1280/1500], Loss: 0.2551
Epoch [1290/1500], Loss: 0.2555
Epoch [1300/1500], Loss: 0.2519
Epoch [1310/1500], Loss: 0.2532
Epoch [1320/1500], Loss: 0.2550
Epoch [1330/1500], Loss: 0.2520
Epoch [1340/1500], Loss: 0.2555
Epoch [1350/1500], Loss: 0.2588
Epoch [1360/1500], Loss: 0.2519
Epoch [1370/1500], Loss: 0.2555
Epoch [1380/1500], Loss: 0.2555
Epoch [1390/1500], Loss: 0.2522
Epoch [1400/1500], Loss: 0.2531
Epoch [1410/1500], Loss: 0.2552
Epoch [1420/1500], Loss: 0.2528
Epoch [1430/1500], Loss: 0.2527
Epoch [1440/1500], Loss: 0.2502
Epoch [1450/1500], Loss: 0.2528
Epoch [1460/1500], Loss: 0.2518
Epoch [1470/1500], Loss: 0.2524
Epoch [1480/1500], Loss: 0.2522
Epoch [1490/1500], Loss: 0.2514
Epoch [1500/1500], Loss: 0.2498
```

Train Accuracy: 0.9372

Test Accuracy: 0.9319

```
In [78]: first_layer_neurons = 256
         dropout_rate = 0.3

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(input_size, 52),
            nn.BatchNorm1d(52),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(52, 26),
            nn.BatchNorm1d(26),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(26, 13),
            nn.BatchNorm1d(13),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            # nn.Linear(13, 6),
            # nn.BatchNorm1d(6),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            # nn.Linear(6, 3),
```

```
        # nn.BatchNorm1d(3),
        # nn.LeakyReLU(0.01),
        # nn.Dropout(dropout_rate),
        nn.Linear(13, 1),
    )

    def forward(self, x):
        return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 800
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/800], Loss: 0.6724  
Epoch [20/800], Loss: 0.6161  
Epoch [30/800], Loss: 0.5721  
Epoch [40/800], Loss: 0.5406  
Epoch [50/800], Loss: 0.5123  
Epoch [60/800], Loss: 0.4936  
Epoch [70/800], Loss: 0.4742  
Epoch [80/800], Loss: 0.4570  
Epoch [90/800], Loss: 0.4417  
Epoch [100/800], Loss: 0.4273  
Epoch [110/800], Loss: 0.4145  
Epoch [120/800], Loss: 0.4034  
Epoch [130/800], Loss: 0.3913  
Epoch [140/800], Loss: 0.3820  
Epoch [150/800], Loss: 0.3718  
Epoch [160/800], Loss: 0.3629  
Epoch [170/800], Loss: 0.3538  
Epoch [180/800], Loss: 0.3467  
Epoch [190/800], Loss: 0.3412  
Epoch [200/800], Loss: 0.3351  
Epoch [210/800], Loss: 0.3274  
Epoch [220/800], Loss: 0.3243  
Epoch [230/800], Loss: 0.3191  
Epoch [240/800], Loss: 0.3157  
Epoch [250/800], Loss: 0.3128  
Epoch [260/800], Loss: 0.3094  
Epoch [270/800], Loss: 0.3032  
Epoch [280/800], Loss: 0.3012  
Epoch [290/800], Loss: 0.2987  
Epoch [300/800], Loss: 0.2958  
Epoch [310/800], Loss: 0.2948  
Epoch [320/800], Loss: 0.2883  
Epoch [330/800], Loss: 0.2848  
Epoch [340/800], Loss: 0.2852  
Epoch [350/800], Loss: 0.2838  
Epoch [360/800], Loss: 0.2791  
Epoch [370/800], Loss: 0.2803  
Epoch [380/800], Loss: 0.2758  
Epoch [390/800], Loss: 0.2737  
Epoch [400/800], Loss: 0.2723  
Epoch [410/800], Loss: 0.2690  
Epoch [420/800], Loss: 0.2696  
Epoch [430/800], Loss: 0.2664  
Epoch [440/800], Loss: 0.2634  
Epoch [450/800], Loss: 0.2624  
Epoch [460/800], Loss: 0.2607  
Epoch [470/800], Loss: 0.2585  
Epoch [480/800], Loss: 0.2579  
Epoch [490/800], Loss: 0.2556  
Epoch [500/800], Loss: 0.2557  
Epoch [510/800], Loss: 0.2533  
Epoch [520/800], Loss: 0.2538  
Epoch [530/800], Loss: 0.2517  
Epoch [540/800], Loss: 0.2455  
Epoch [550/800], Loss: 0.2504  
Epoch [560/800], Loss: 0.2463  
Epoch [570/800], Loss: 0.2450  
Epoch [580/800], Loss: 0.2446  
Epoch [590/800], Loss: 0.2449  
Epoch [600/800], Loss: 0.2446

```
Epoch [610/800], Loss: 0.2420
Epoch [620/800], Loss: 0.2403
Epoch [630/800], Loss: 0.2396
Epoch [640/800], Loss: 0.2360
Epoch [650/800], Loss: 0.2401
Epoch [660/800], Loss: 0.2357
Epoch [670/800], Loss: 0.2335
Epoch [680/800], Loss: 0.2345
Epoch [690/800], Loss: 0.2319
Epoch [700/800], Loss: 0.2344
Epoch [710/800], Loss: 0.2311
Epoch [720/800], Loss: 0.2328
Epoch [730/800], Loss: 0.2327
Epoch [740/800], Loss: 0.2320
Epoch [750/800], Loss: 0.2317
Epoch [760/800], Loss: 0.2290
Epoch [770/800], Loss: 0.2306
Epoch [780/800], Loss: 0.2285
Epoch [790/800], Loss: 0.2296
Epoch [800/800], Loss: 0.2272
```

Train Accuracy: 0.9324

Test Accuracy: 0.9289

```
In [79]: first_layer_neurons = 256
         dropout_rate = 0.3

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(input_size, 104),
            nn.BatchNorm1d(104),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(104, 52),
            nn.BatchNorm1d(52),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(52, 26),
            nn.BatchNorm1d(26),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            # nn.Linear(26, 13),
            # nn.BatchNorm1d(13),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            # nn.Linear(13, 6),
            # nn.BatchNorm1d(6),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            # nn.Linear(6, 3),
            # nn.BatchNorm1d(3),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            nn.Linear(26, 1),
        )
```

```
def forward(self, x):
    return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 2000
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/2000], Loss: 0.6166  
Epoch [20/2000], Loss: 0.5379  
Epoch [30/2000], Loss: 0.4846  
Epoch [40/2000], Loss: 0.4485  
Epoch [50/2000], Loss: 0.4231  
Epoch [60/2000], Loss: 0.3999  
Epoch [70/2000], Loss: 0.3807  
Epoch [80/2000], Loss: 0.3657  
Epoch [90/2000], Loss: 0.3506  
Epoch [100/2000], Loss: 0.3408  
Epoch [110/2000], Loss: 0.3321  
Epoch [120/2000], Loss: 0.3200  
Epoch [130/2000], Loss: 0.3133  
Epoch [140/2000], Loss: 0.3088  
Epoch [150/2000], Loss: 0.3015  
Epoch [160/2000], Loss: 0.2961  
Epoch [170/2000], Loss: 0.2918  
Epoch [180/2000], Loss: 0.2855  
Epoch [190/2000], Loss: 0.2821  
Epoch [200/2000], Loss: 0.2777  
Epoch [210/2000], Loss: 0.2759  
Epoch [220/2000], Loss: 0.2753  
Epoch [230/2000], Loss: 0.2699  
Epoch [240/2000], Loss: 0.2704  
Epoch [250/2000], Loss: 0.2634  
Epoch [260/2000], Loss: 0.2599  
Epoch [270/2000], Loss: 0.2606  
Epoch [280/2000], Loss: 0.2561  
Epoch [290/2000], Loss: 0.2540  
Epoch [300/2000], Loss: 0.2526  
Epoch [310/2000], Loss: 0.2513  
Epoch [320/2000], Loss: 0.2470  
Epoch [330/2000], Loss: 0.2468  
Epoch [340/2000], Loss: 0.2462  
Epoch [350/2000], Loss: 0.2464  
Epoch [360/2000], Loss: 0.2417  
Epoch [370/2000], Loss: 0.2419  
Epoch [380/2000], Loss: 0.2395  
Epoch [390/2000], Loss: 0.2384  
Epoch [400/2000], Loss: 0.2376  
Epoch [410/2000], Loss: 0.2319  
Epoch [420/2000], Loss: 0.2325  
Epoch [430/2000], Loss: 0.2319  
Epoch [440/2000], Loss: 0.2285  
Epoch [450/2000], Loss: 0.2299  
Epoch [460/2000], Loss: 0.2251  
Epoch [470/2000], Loss: 0.2233  
Epoch [480/2000], Loss: 0.2248  
Epoch [490/2000], Loss: 0.2226  
Epoch [500/2000], Loss: 0.2209  
Epoch [510/2000], Loss: 0.2169  
Epoch [520/2000], Loss: 0.2207  
Epoch [530/2000], Loss: 0.2164  
Epoch [540/2000], Loss: 0.2158  
Epoch [550/2000], Loss: 0.2153  
Epoch [560/2000], Loss: 0.2153  
Epoch [570/2000], Loss: 0.2137  
Epoch [580/2000], Loss: 0.2144  
Epoch [590/2000], Loss: 0.2118  
Epoch [600/2000], Loss: 0.2102

Epoch [610/2000], Loss: 0.2109  
Epoch [620/2000], Loss: 0.2070  
Epoch [630/2000], Loss: 0.2076  
Epoch [640/2000], Loss: 0.2087  
Epoch [650/2000], Loss: 0.2078  
Epoch [660/2000], Loss: 0.2072  
Epoch [670/2000], Loss: 0.2066  
Epoch [680/2000], Loss: 0.2064  
Epoch [690/2000], Loss: 0.2069  
Epoch [700/2000], Loss: 0.2033  
Epoch [710/2000], Loss: 0.2025  
Epoch [720/2000], Loss: 0.2021  
Epoch [730/2000], Loss: 0.2014  
Epoch [740/2000], Loss: 0.2037  
Epoch [750/2000], Loss: 0.2032  
Epoch [760/2000], Loss: 0.2015  
Epoch [770/2000], Loss: 0.2027  
Epoch [780/2000], Loss: 0.2003  
Epoch [790/2000], Loss: 0.2018  
Epoch [800/2000], Loss: 0.2025  
Epoch [810/2000], Loss: 0.1999  
Epoch [820/2000], Loss: 0.2004  
Epoch [830/2000], Loss: 0.1988  
Epoch [840/2000], Loss: 0.1950  
Epoch [850/2000], Loss: 0.1997  
Epoch [860/2000], Loss: 0.1977  
Epoch [870/2000], Loss: 0.1986  
Epoch [880/2000], Loss: 0.1967  
Epoch [890/2000], Loss: 0.1976  
Epoch [900/2000], Loss: 0.1970  
Epoch [910/2000], Loss: 0.1985  
Epoch [920/2000], Loss: 0.1967  
Epoch [930/2000], Loss: 0.1940  
Epoch [940/2000], Loss: 0.1955  
Epoch [950/2000], Loss: 0.1945  
Epoch [960/2000], Loss: 0.1945  
Epoch [970/2000], Loss: 0.1919  
Epoch [980/2000], Loss: 0.1946  
Epoch [990/2000], Loss: 0.1958  
Epoch [1000/2000], Loss: 0.1931  
Epoch [1010/2000], Loss: 0.1948  
Epoch [1020/2000], Loss: 0.1898  
Epoch [1030/2000], Loss: 0.1939  
Epoch [1040/2000], Loss: 0.1933  
Epoch [1050/2000], Loss: 0.1936  
Epoch [1060/2000], Loss: 0.1936  
Epoch [1070/2000], Loss: 0.1958  
Epoch [1080/2000], Loss: 0.1921  
Epoch [1090/2000], Loss: 0.1921  
Epoch [1100/2000], Loss: 0.1914  
Epoch [1110/2000], Loss: 0.1929  
Epoch [1120/2000], Loss: 0.1894  
Epoch [1130/2000], Loss: 0.1918  
Epoch [1140/2000], Loss: 0.1944  
Epoch [1150/2000], Loss: 0.1949  
Epoch [1160/2000], Loss: 0.1903  
Epoch [1170/2000], Loss: 0.1871  
Epoch [1180/2000], Loss: 0.1894  
Epoch [1190/2000], Loss: 0.1913  
Epoch [1200/2000], Loss: 0.1879

Epoch [1210/2000], Loss: 0.1910  
Epoch [1220/2000], Loss: 0.1909  
Epoch [1230/2000], Loss: 0.1863  
Epoch [1240/2000], Loss: 0.1900  
Epoch [1250/2000], Loss: 0.1867  
Epoch [1260/2000], Loss: 0.1863  
Epoch [1270/2000], Loss: 0.1874  
Epoch [1280/2000], Loss: 0.1896  
Epoch [1290/2000], Loss: 0.1882  
Epoch [1300/2000], Loss: 0.1883  
Epoch [1310/2000], Loss: 0.1881  
Epoch [1320/2000], Loss: 0.1878  
Epoch [1330/2000], Loss: 0.1869  
Epoch [1340/2000], Loss: 0.1879  
Epoch [1350/2000], Loss: 0.1863  
Epoch [1360/2000], Loss: 0.1874  
Epoch [1370/2000], Loss: 0.1902  
Epoch [1380/2000], Loss: 0.1873  
Epoch [1390/2000], Loss: 0.1842  
Epoch [1400/2000], Loss: 0.1865  
Epoch [1410/2000], Loss: 0.1850  
Epoch [1420/2000], Loss: 0.1858  
Epoch [1430/2000], Loss: 0.1863  
Epoch [1440/2000], Loss: 0.1842  
Epoch [1450/2000], Loss: 0.1860  
Epoch [1460/2000], Loss: 0.1852  
Epoch [1470/2000], Loss: 0.1856  
Epoch [1480/2000], Loss: 0.1845  
Epoch [1490/2000], Loss: 0.1878  
Epoch [1500/2000], Loss: 0.1888  
Epoch [1510/2000], Loss: 0.1853  
Epoch [1520/2000], Loss: 0.1844  
Epoch [1530/2000], Loss: 0.1844  
Epoch [1540/2000], Loss: 0.1863  
Epoch [1550/2000], Loss: 0.1855  
Epoch [1560/2000], Loss: 0.1833  
Epoch [1570/2000], Loss: 0.1865  
Epoch [1580/2000], Loss: 0.1846  
Epoch [1590/2000], Loss: 0.1841  
Epoch [1600/2000], Loss: 0.1866  
Epoch [1610/2000], Loss: 0.1848  
Epoch [1620/2000], Loss: 0.1841  
Epoch [1630/2000], Loss: 0.1853  
Epoch [1640/2000], Loss: 0.1871  
Epoch [1650/2000], Loss: 0.1831  
Epoch [1660/2000], Loss: 0.1863  
Epoch [1670/2000], Loss: 0.1857  
Epoch [1680/2000], Loss: 0.1857  
Epoch [1690/2000], Loss: 0.1832  
Epoch [1700/2000], Loss: 0.1830  
Epoch [1710/2000], Loss: 0.1819  
Epoch [1720/2000], Loss: 0.1822  
Epoch [1730/2000], Loss: 0.1800  
Epoch [1740/2000], Loss: 0.1827  
Epoch [1750/2000], Loss: 0.1862  
Epoch [1760/2000], Loss: 0.1848  
Epoch [1770/2000], Loss: 0.1829  
Epoch [1780/2000], Loss: 0.1844  
Epoch [1790/2000], Loss: 0.1853  
Epoch [1800/2000], Loss: 0.1830



```
Epoch [1810/2000], Loss: 0.1830
Epoch [1820/2000], Loss: 0.1822
Epoch [1830/2000], Loss: 0.1815
Epoch [1840/2000], Loss: 0.1831
Epoch [1850/2000], Loss: 0.1841
Epoch [1860/2000], Loss: 0.1840
Epoch [1870/2000], Loss: 0.1824
Epoch [1880/2000], Loss: 0.1842
Epoch [1890/2000], Loss: 0.1811
Epoch [1900/2000], Loss: 0.1825
Epoch [1910/2000], Loss: 0.1801
Epoch [1920/2000], Loss: 0.1844
Epoch [1930/2000], Loss: 0.1824
Epoch [1940/2000], Loss: 0.1811
Epoch [1950/2000], Loss: 0.1852
Epoch [1960/2000], Loss: 0.1807
Epoch [1970/2000], Loss: 0.1807
Epoch [1980/2000], Loss: 0.1813
Epoch [1990/2000], Loss: 0.1827
Epoch [2000/2000], Loss: 0.1835
```

Train Accuracy: 0.9404

Test Accuracy: 0.9282

```
In [80]: first_layer_neurons = 256
         dropout_rate = 0.3

class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(input_size, 256),
            nn.BatchNorm1d(256),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(256, 64),
            nn.BatchNorm1d(64),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(64, 16),
            nn.BatchNorm1d(16),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            # nn.Linear(26, 13),
            # nn.BatchNorm1d(13),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            # nn.Linear(13, 6),
            # nn.BatchNorm1d(6),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            # nn.Linear(6, 3),
            # nn.BatchNorm1d(3),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            nn.Linear(16, 1),
        )
```

```
def forward(self, x):
    return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 2000
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/2000], Loss: 0.5050  
Epoch [20/2000], Loss: 0.4627  
Epoch [30/2000], Loss: 0.4346  
Epoch [40/2000], Loss: 0.4147  
Epoch [50/2000], Loss: 0.3965  
Epoch [60/2000], Loss: 0.3795  
Epoch [70/2000], Loss: 0.3667  
Epoch [80/2000], Loss: 0.3521  
Epoch [90/2000], Loss: 0.3400  
Epoch [100/2000], Loss: 0.3303  
Epoch [110/2000], Loss: 0.3196  
Epoch [120/2000], Loss: 0.3096  
Epoch [130/2000], Loss: 0.3046  
Epoch [140/2000], Loss: 0.2967  
Epoch [150/2000], Loss: 0.2889  
Epoch [160/2000], Loss: 0.2854  
Epoch [170/2000], Loss: 0.2791  
Epoch [180/2000], Loss: 0.2745  
Epoch [190/2000], Loss: 0.2719  
Epoch [200/2000], Loss: 0.2683  
Epoch [210/2000], Loss: 0.2622  
Epoch [220/2000], Loss: 0.2584  
Epoch [230/2000], Loss: 0.2562  
Epoch [240/2000], Loss: 0.2531  
Epoch [250/2000], Loss: 0.2485  
Epoch [260/2000], Loss: 0.2489  
Epoch [270/2000], Loss: 0.2458  
Epoch [280/2000], Loss: 0.2436  
Epoch [290/2000], Loss: 0.2425  
Epoch [300/2000], Loss: 0.2416  
Epoch [310/2000], Loss: 0.2355  
Epoch [320/2000], Loss: 0.2357  
Epoch [330/2000], Loss: 0.2350  
Epoch [340/2000], Loss: 0.2312  
Epoch [350/2000], Loss: 0.2298  
Epoch [360/2000], Loss: 0.2275  
Epoch [370/2000], Loss: 0.2275  
Epoch [380/2000], Loss: 0.2283  
Epoch [390/2000], Loss: 0.2257  
Epoch [400/2000], Loss: 0.2211  
Epoch [410/2000], Loss: 0.2192  
Epoch [420/2000], Loss: 0.2176  
Epoch [430/2000], Loss: 0.2188  
Epoch [440/2000], Loss: 0.2170  
Epoch [450/2000], Loss: 0.2148  
Epoch [460/2000], Loss: 0.2129  
Epoch [470/2000], Loss: 0.2136  
Epoch [480/2000], Loss: 0.2104  
Epoch [490/2000], Loss: 0.2099  
Epoch [500/2000], Loss: 0.2084  
Epoch [510/2000], Loss: 0.2060  
Epoch [520/2000], Loss: 0.2080  
Epoch [530/2000], Loss: 0.2065  
Epoch [540/2000], Loss: 0.2034  
Epoch [550/2000], Loss: 0.2020  
Epoch [560/2000], Loss: 0.2047  
Epoch [570/2000], Loss: 0.1995  
Epoch [580/2000], Loss: 0.1971  
Epoch [590/2000], Loss: 0.1960  
Epoch [600/2000], Loss: 0.1975

Epoch [610/2000], Loss: 0.1955  
Epoch [620/2000], Loss: 0.1940  
Epoch [630/2000], Loss: 0.1924  
Epoch [640/2000], Loss: 0.1942  
Epoch [650/2000], Loss: 0.1916  
Epoch [660/2000], Loss: 0.1937  
Epoch [670/2000], Loss: 0.1926  
Epoch [680/2000], Loss: 0.1886  
Epoch [690/2000], Loss: 0.1912  
Epoch [700/2000], Loss: 0.1878  
Epoch [710/2000], Loss: 0.1868  
Epoch [720/2000], Loss: 0.1856  
Epoch [730/2000], Loss: 0.1881  
Epoch [740/2000], Loss: 0.1837  
Epoch [750/2000], Loss: 0.1847  
Epoch [760/2000], Loss: 0.1845  
Epoch [770/2000], Loss: 0.1808  
Epoch [780/2000], Loss: 0.1839  
Epoch [790/2000], Loss: 0.1833  
Epoch [800/2000], Loss: 0.1826  
Epoch [810/2000], Loss: 0.1811  
Epoch [820/2000], Loss: 0.1816  
Epoch [830/2000], Loss: 0.1806  
Epoch [840/2000], Loss: 0.1794  
Epoch [850/2000], Loss: 0.1768  
Epoch [860/2000], Loss: 0.1748  
Epoch [870/2000], Loss: 0.1775  
Epoch [880/2000], Loss: 0.1754  
Epoch [890/2000], Loss: 0.1773  
Epoch [900/2000], Loss: 0.1752  
Epoch [910/2000], Loss: 0.1766  
Epoch [920/2000], Loss: 0.1746  
Epoch [930/2000], Loss: 0.1724  
Epoch [940/2000], Loss: 0.1751  
Epoch [950/2000], Loss: 0.1697  
Epoch [960/2000], Loss: 0.1718  
Epoch [970/2000], Loss: 0.1693  
Epoch [980/2000], Loss: 0.1688  
Epoch [990/2000], Loss: 0.1679  
Epoch [1000/2000], Loss: 0.1705  
Epoch [1010/2000], Loss: 0.1678  
Epoch [1020/2000], Loss: 0.1675  
Epoch [1030/2000], Loss: 0.1680  
Epoch [1040/2000], Loss: 0.1675  
Epoch [1050/2000], Loss: 0.1674  
Epoch [1060/2000], Loss: 0.1659  
Epoch [1070/2000], Loss: 0.1672  
Epoch [1080/2000], Loss: 0.1664  
Epoch [1090/2000], Loss: 0.1613  
Epoch [1100/2000], Loss: 0.1618  
Epoch [1110/2000], Loss: 0.1643  
Epoch [1120/2000], Loss: 0.1634  
Epoch [1130/2000], Loss: 0.1638  
Epoch [1140/2000], Loss: 0.1632  
Epoch [1150/2000], Loss: 0.1599  
Epoch [1160/2000], Loss: 0.1634  
Epoch [1170/2000], Loss: 0.1591  
Epoch [1180/2000], Loss: 0.1592  
Epoch [1190/2000], Loss: 0.1599  
Epoch [1200/2000], Loss: 0.1611

Epoch [1210/2000], Loss: 0.1593  
Epoch [1220/2000], Loss: 0.1592  
Epoch [1230/2000], Loss: 0.1621  
Epoch [1240/2000], Loss: 0.1584  
Epoch [1250/2000], Loss: 0.1599  
Epoch [1260/2000], Loss: 0.1549  
Epoch [1270/2000], Loss: 0.1588  
Epoch [1280/2000], Loss: 0.1567  
Epoch [1290/2000], Loss: 0.1585  
Epoch [1300/2000], Loss: 0.1551  
Epoch [1310/2000], Loss: 0.1590  
Epoch [1320/2000], Loss: 0.1591  
Epoch [1330/2000], Loss: 0.1552  
Epoch [1340/2000], Loss: 0.1546  
Epoch [1350/2000], Loss: 0.1534  
Epoch [1360/2000], Loss: 0.1547  
Epoch [1370/2000], Loss: 0.1545  
Epoch [1380/2000], Loss: 0.1522  
Epoch [1390/2000], Loss: 0.1520  
Epoch [1400/2000], Loss: 0.1557  
Epoch [1410/2000], Loss: 0.1535  
Epoch [1420/2000], Loss: 0.1503  
Epoch [1430/2000], Loss: 0.1502  
Epoch [1440/2000], Loss: 0.1553  
Epoch [1450/2000], Loss: 0.1542  
Epoch [1460/2000], Loss: 0.1528  
Epoch [1470/2000], Loss: 0.1513  
Epoch [1480/2000], Loss: 0.1547  
Epoch [1490/2000], Loss: 0.1527  
Epoch [1500/2000], Loss: 0.1528  
Epoch [1510/2000], Loss: 0.1505  
Epoch [1520/2000], Loss: 0.1533  
Epoch [1530/2000], Loss: 0.1506  
Epoch [1540/2000], Loss: 0.1503  
Epoch [1550/2000], Loss: 0.1474  
Epoch [1560/2000], Loss: 0.1493  
Epoch [1570/2000], Loss: 0.1496  
Epoch [1580/2000], Loss: 0.1507  
Epoch [1590/2000], Loss: 0.1487  
Epoch [1600/2000], Loss: 0.1492  
Epoch [1610/2000], Loss: 0.1497  
Epoch [1620/2000], Loss: 0.1456  
Epoch [1630/2000], Loss: 0.1456  
Epoch [1640/2000], Loss: 0.1448  
Epoch [1650/2000], Loss: 0.1486  
Epoch [1660/2000], Loss: 0.1453  
Epoch [1670/2000], Loss: 0.1476  
Epoch [1680/2000], Loss: 0.1455  
Epoch [1690/2000], Loss: 0.1488  
Epoch [1700/2000], Loss: 0.1481  
Epoch [1710/2000], Loss: 0.1473  
Epoch [1720/2000], Loss: 0.1444  
Epoch [1730/2000], Loss: 0.1479  
Epoch [1740/2000], Loss: 0.1502  
Epoch [1750/2000], Loss: 0.1467  
Epoch [1760/2000], Loss: 0.1429  
Epoch [1770/2000], Loss: 0.1461  
Epoch [1780/2000], Loss: 0.1421  
Epoch [1790/2000], Loss: 0.1432  
Epoch [1800/2000], Loss: 0.1508

```
Epoch [1810/2000], Loss: 0.1433
Epoch [1820/2000], Loss: 0.1461
Epoch [1830/2000], Loss: 0.1452
Epoch [1840/2000], Loss: 0.1439
Epoch [1850/2000], Loss: 0.1466
Epoch [1860/2000], Loss: 0.1456
Epoch [1870/2000], Loss: 0.1428
Epoch [1880/2000], Loss: 0.1457
Epoch [1890/2000], Loss: 0.1448
Epoch [1900/2000], Loss: 0.1438
Epoch [1910/2000], Loss: 0.1460
Epoch [1920/2000], Loss: 0.1439
Epoch [1930/2000], Loss: 0.1400
Epoch [1940/2000], Loss: 0.1429
Epoch [1950/2000], Loss: 0.1440
Epoch [1960/2000], Loss: 0.1423
Epoch [1970/2000], Loss: 0.1412
Epoch [1980/2000], Loss: 0.1415
Epoch [1990/2000], Loss: 0.1414
Epoch [2000/2000], Loss: 0.1427
```

Train Accuracy: 0.9559

Test Accuracy: 0.9281

In [81]: dropout\_rate = 0.4

```
class LoanDefaultNN(nn.Module):
    def __init__(self, input_size):
        super(LoanDefaultNN, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(input_size, 104),
            nn.BatchNorm1d(104),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(104, 52),
            nn.BatchNorm1d(52),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            nn.Linear(52, 26),
            nn.BatchNorm1d(26),
            nn.LeakyReLU(0.01),
            nn.Dropout(dropout_rate),
            # nn.Linear(26, 13),
            # nn.BatchNorm1d(13),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            # nn.Linear(13, 6),
            # nn.BatchNorm1d(6),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            # nn.Linear(6, 3),
            # nn.BatchNorm1d(3),
            # nn.LeakyReLU(0.01),
            # nn.Dropout(dropout_rate),
            nn.Linear(26, 1),
        )

    def forward(self, x):
```

```
        return self.model(x) # Use if you're sticking with BCELoss

# Initialize model, loss, optimizer
model = LoanDefaultNN(X_train.shape[1])
criterion = nn.BCEWithLogitsLoss()
# Add regularization in weight decay
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-3)

# Training loop
epochs = 2000
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    output = model(X_train_tensor)
    loss = criterion(output, y_train_tensor)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")

model.eval()
with torch.no_grad():
    test_preds = model(X_train_tensor)
    test_predicted_classes = (test_preds > 0.5).float()
    test_accuracy = (test_predicted_classes == y_train_tensor).float().mean()
    print(f"\nTrain Accuracy: {test_accuracy.item():.4f}")
    preds = model(X_test_tensor)
    predicted_classes = (preds > 0.5).float()
    accuracy = (predicted_classes == y_test_tensor).float().mean()
    print(f"\nTest Accuracy: {accuracy.item():.4f}")
```

Epoch [10/2000], Loss: 0.6861  
Epoch [20/2000], Loss: 0.6102  
Epoch [30/2000], Loss: 0.5577  
Epoch [40/2000], Loss: 0.5176  
Epoch [50/2000], Loss: 0.4910  
Epoch [60/2000], Loss: 0.4661  
Epoch [70/2000], Loss: 0.4457  
Epoch [80/2000], Loss: 0.4286  
Epoch [90/2000], Loss: 0.4122  
Epoch [100/2000], Loss: 0.3970  
Epoch [110/2000], Loss: 0.3861  
Epoch [120/2000], Loss: 0.3719  
Epoch [130/2000], Loss: 0.3641  
Epoch [140/2000], Loss: 0.3505  
Epoch [150/2000], Loss: 0.3449  
Epoch [160/2000], Loss: 0.3349  
Epoch [170/2000], Loss: 0.3253  
Epoch [180/2000], Loss: 0.3196  
Epoch [190/2000], Loss: 0.3138  
Epoch [200/2000], Loss: 0.3077  
Epoch [210/2000], Loss: 0.3032  
Epoch [220/2000], Loss: 0.2978  
Epoch [230/2000], Loss: 0.2936  
Epoch [240/2000], Loss: 0.2913  
Epoch [250/2000], Loss: 0.2866  
Epoch [260/2000], Loss: 0.2831  
Epoch [270/2000], Loss: 0.2821  
Epoch [280/2000], Loss: 0.2776  
Epoch [290/2000], Loss: 0.2756  
Epoch [300/2000], Loss: 0.2716  
Epoch [310/2000], Loss: 0.2722  
Epoch [320/2000], Loss: 0.2716  
Epoch [330/2000], Loss: 0.2636  
Epoch [340/2000], Loss: 0.2660  
Epoch [350/2000], Loss: 0.2614  
Epoch [360/2000], Loss: 0.2594  
Epoch [370/2000], Loss: 0.2574  
Epoch [380/2000], Loss: 0.2548  
Epoch [390/2000], Loss: 0.2525  
Epoch [400/2000], Loss: 0.2531  
Epoch [410/2000], Loss: 0.2505  
Epoch [420/2000], Loss: 0.2468  
Epoch [430/2000], Loss: 0.2470  
Epoch [440/2000], Loss: 0.2406  
Epoch [450/2000], Loss: 0.2417  
Epoch [460/2000], Loss: 0.2380  
Epoch [470/2000], Loss: 0.2387  
Epoch [480/2000], Loss: 0.2389  
Epoch [490/2000], Loss: 0.2369  
Epoch [500/2000], Loss: 0.2343  
Epoch [510/2000], Loss: 0.2289  
Epoch [520/2000], Loss: 0.2316  
Epoch [530/2000], Loss: 0.2293  
Epoch [540/2000], Loss: 0.2295  
Epoch [550/2000], Loss: 0.2305  
Epoch [560/2000], Loss: 0.2286  
Epoch [570/2000], Loss: 0.2286  
Epoch [580/2000], Loss: 0.2247  
Epoch [590/2000], Loss: 0.2240  
Epoch [600/2000], Loss: 0.2257



Epoch [610/2000], Loss: 0.2259  
Epoch [620/2000], Loss: 0.2231  
Epoch [630/2000], Loss: 0.2237  
Epoch [640/2000], Loss: 0.2217  
Epoch [650/2000], Loss: 0.2240  
Epoch [660/2000], Loss: 0.2226  
Epoch [670/2000], Loss: 0.2199  
Epoch [680/2000], Loss: 0.2197  
Epoch [690/2000], Loss: 0.2195  
Epoch [700/2000], Loss: 0.2189  
Epoch [710/2000], Loss: 0.2207  
Epoch [720/2000], Loss: 0.2201  
Epoch [730/2000], Loss: 0.2196  
Epoch [740/2000], Loss: 0.2201  
Epoch [750/2000], Loss: 0.2149  
Epoch [760/2000], Loss: 0.2158  
Epoch [770/2000], Loss: 0.2154  
Epoch [780/2000], Loss: 0.2149  
Epoch [790/2000], Loss: 0.2135  
Epoch [800/2000], Loss: 0.2143  
Epoch [810/2000], Loss: 0.2138  
Epoch [820/2000], Loss: 0.2173  
Epoch [830/2000], Loss: 0.2152  
Epoch [840/2000], Loss: 0.2137  
Epoch [850/2000], Loss: 0.2126  
Epoch [860/2000], Loss: 0.2105  
Epoch [870/2000], Loss: 0.2133  
Epoch [880/2000], Loss: 0.2094  
Epoch [890/2000], Loss: 0.2116  
Epoch [900/2000], Loss: 0.2138  
Epoch [910/2000], Loss: 0.2115  
Epoch [920/2000], Loss: 0.2091  
Epoch [930/2000], Loss: 0.2111  
Epoch [940/2000], Loss: 0.2107  
Epoch [950/2000], Loss: 0.2109  
Epoch [960/2000], Loss: 0.2107  
Epoch [970/2000], Loss: 0.2086  
Epoch [980/2000], Loss: 0.2102  
Epoch [990/2000], Loss: 0.2101  
Epoch [1000/2000], Loss: 0.2071  
Epoch [1010/2000], Loss: 0.2105  
Epoch [1020/2000], Loss: 0.2088  
Epoch [1030/2000], Loss: 0.2132  
Epoch [1040/2000], Loss: 0.2096  
Epoch [1050/2000], Loss: 0.2088  
Epoch [1060/2000], Loss: 0.2103  
Epoch [1070/2000], Loss: 0.2083  
Epoch [1080/2000], Loss: 0.2101  
Epoch [1090/2000], Loss: 0.2073  
Epoch [1100/2000], Loss: 0.2085  
Epoch [1110/2000], Loss: 0.2064  
Epoch [1120/2000], Loss: 0.2067  
Epoch [1130/2000], Loss: 0.2088  
Epoch [1140/2000], Loss: 0.2066  
Epoch [1150/2000], Loss: 0.2073  
Epoch [1160/2000], Loss: 0.2052  
Epoch [1170/2000], Loss: 0.2063  
Epoch [1180/2000], Loss: 0.2051  
Epoch [1190/2000], Loss: 0.2071  
Epoch [1200/2000], Loss: 0.2063

Epoch [1210/2000], Loss: 0.2074  
Epoch [1220/2000], Loss: 0.2054  
Epoch [1230/2000], Loss: 0.2067  
Epoch [1240/2000], Loss: 0.2057  
Epoch [1250/2000], Loss: 0.2050  
Epoch [1260/2000], Loss: 0.2065  
Epoch [1270/2000], Loss: 0.2046  
Epoch [1280/2000], Loss: 0.2045  
Epoch [1290/2000], Loss: 0.2038  
Epoch [1300/2000], Loss: 0.2067  
Epoch [1310/2000], Loss: 0.2051  
Epoch [1320/2000], Loss: 0.2055  
Epoch [1330/2000], Loss: 0.2042  
Epoch [1340/2000], Loss: 0.2034  
Epoch [1350/2000], Loss: 0.2049  
Epoch [1360/2000], Loss: 0.2057  
Epoch [1370/2000], Loss: 0.2044  
Epoch [1380/2000], Loss: 0.2040  
Epoch [1390/2000], Loss: 0.2028  
Epoch [1400/2000], Loss: 0.2041  
Epoch [1410/2000], Loss: 0.2048  
Epoch [1420/2000], Loss: 0.2028  
Epoch [1430/2000], Loss: 0.2033  
Epoch [1440/2000], Loss: 0.2038  
Epoch [1450/2000], Loss: 0.2037  
Epoch [1460/2000], Loss: 0.2036  
Epoch [1470/2000], Loss: 0.2007  
Epoch [1480/2000], Loss: 0.2023  
Epoch [1490/2000], Loss: 0.2031  
Epoch [1500/2000], Loss: 0.2027  
Epoch [1510/2000], Loss: 0.2029  
Epoch [1520/2000], Loss: 0.2012  
Epoch [1530/2000], Loss: 0.2025  
Epoch [1540/2000], Loss: 0.2039  
Epoch [1550/2000], Loss: 0.2010  
Epoch [1560/2000], Loss: 0.2035  
Epoch [1570/2000], Loss: 0.2010  
Epoch [1580/2000], Loss: 0.2020  
Epoch [1590/2000], Loss: 0.2050  
Epoch [1600/2000], Loss: 0.2040  
Epoch [1610/2000], Loss: 0.2024  
Epoch [1620/2000], Loss: 0.2054  
Epoch [1630/2000], Loss: 0.2005  
Epoch [1640/2000], Loss: 0.2041  
Epoch [1650/2000], Loss: 0.2033  
Epoch [1660/2000], Loss: 0.2011  
Epoch [1670/2000], Loss: 0.1986  
Epoch [1680/2000], Loss: 0.2022  
Epoch [1690/2000], Loss: 0.2016  
Epoch [1700/2000], Loss: 0.2014  
Epoch [1710/2000], Loss: 0.2018  
Epoch [1720/2000], Loss: 0.2022  
Epoch [1730/2000], Loss: 0.2024  
Epoch [1740/2000], Loss: 0.2024  
Epoch [1750/2000], Loss: 0.2031  
Epoch [1760/2000], Loss: 0.1989  
Epoch [1770/2000], Loss: 0.1983  
Epoch [1780/2000], Loss: 0.2004  
Epoch [1790/2000], Loss: 0.1994  
Epoch [1800/2000], Loss: 0.2009

```
Epoch [1810/2000], Loss: 0.2024
Epoch [1820/2000], Loss: 0.2020
Epoch [1830/2000], Loss: 0.2008
Epoch [1840/2000], Loss: 0.1977
Epoch [1850/2000], Loss: 0.2001
Epoch [1860/2000], Loss: 0.2011
Epoch [1870/2000], Loss: 0.1998
Epoch [1880/2000], Loss: 0.2053
Epoch [1890/2000], Loss: 0.2009
Epoch [1900/2000], Loss: 0.2022
Epoch [1910/2000], Loss: 0.1985
Epoch [1920/2000], Loss: 0.2019
Epoch [1930/2000], Loss: 0.2001
Epoch [1940/2000], Loss: 0.2008
Epoch [1950/2000], Loss: 0.1982
Epoch [1960/2000], Loss: 0.2012
Epoch [1970/2000], Loss: 0.1999
Epoch [1980/2000], Loss: 0.1993
Epoch [1990/2000], Loss: 0.1994
Epoch [2000/2000], Loss: 0.2001
```

Train Accuracy: 0.9384

Test Accuracy: 0.9305

## ML Method 2: Decision Trees:

Decision trees are highly effective for our use case because they offer clear interpretability.

Unlike neural networks, which often operate as "black boxes" with complex internal structures, decision trees provide greater observability and traceability.

We tested multiple tree types : single multi-variable trees, random forest, and XGBOOST. and consolidated all their results and code into a single cell below

```
In [82]: df = df.dropna()
X = df.drop("loan_status", axis=1)
y = df["loan_status"]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=1015
)
```

```
In [83]: import pandas as pd
import xgboost as xgb
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    confusion_matrix,
```

```
classification_report,  
)  
  
# Create a function to train, evaluate and visualize models  
def evaluate_model(model, model_name, X_train, X_test, y_train, y_test):  
    # Train the model  
    model.fit(X_train, y_train)  
  
    # Make predictions on train and test sets  
    y_train_pred = model.predict(X_train)  
    y_test_pred = model.predict(X_test)  
  
    # Calculate metrics for training data  
    train_accuracy = accuracy_score(y_train, y_train_pred)  
    train_precision = precision_score(y_train, y_train_pred)  
    train_recall = recall_score(y_train, y_train_pred)  
    train_f1 = f1_score(y_train, y_train_pred)  
  
    # Calculate metrics for test data  
    test_accuracy = accuracy_score(y_test, y_test_pred)  
    test_precision = precision_score(y_test, y_test_pred)  
    test_recall = recall_score(y_test, y_test_pred)  
    test_f1 = f1_score(y_test, y_test_pred)  
  
    # Print results  
    print(f"\n--- {model_name} Results ---")  
    print(f"Training Metrics:")  
    print(f"Accuracy: {train_accuracy:.4f}")  
    print(f"Precision: {train_precision:.4f}")  
    print(f"Recall: {train_recall:.4f}")  
    print(f"F1 Score: {train_f1:.4f}")  
  
    print(f"\nTest Metrics:")  
    print(f"Accuracy: {test_accuracy:.4f}")  
    print(f"Precision: {test_precision:.4f}")  
    print(f"Recall: {test_recall:.4f}")  
    print(f"F1 Score: {test_f1:.4f}")  
  
    # Print classification reports  
    print(f"\nTraining Classification Report:")  
    print(classification_report(y_train, y_train_pred))  
  
    print(f"\nTest Classification Report:")  
    print(classification_report(y_test, y_test_pred))  
  
    # Calculate confusion matrices  
    train_cm = confusion_matrix(y_train, y_train_pred)  
    test_cm = confusion_matrix(y_test, y_test_pred)  
  
    # Print confusion matrices  
    print(f"\nTraining Confusion Matrix:")  
    print(train_cm)  
  
    print(f"\nTest Confusion Matrix:")  
    print(test_cm)  
  
    # Plot confusion matrices  
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))
```

```
# Training confusion matrix
sns.heatmap(
    train_cm,
    annot=True,
    fmt="d",
    cmap="Blues",
    xticklabels=["0", "1"],
    yticklabels=["0", "1"],
    ax=ax1,
)
ax1.set_xlabel("Predicted")
ax1.set_ylabel("Actual")
ax1.set_title(f"{model_name} - Training Confusion Matrix")

# Test confusion matrix
sns.heatmap(
    test_cm,
    annot=True,
    fmt="d",
    cmap="Blues",
    xticklabels=["0", "1"],
    yticklabels=["0", "1"],
    ax=ax2,
)
ax2.set_xlabel("Predicted")
ax2.set_ylabel("Actual")
ax2.set_title(f"{model_name} - Test Confusion Matrix")

plt.tight_layout()
plt.show()

# Initialize the models
decision_tree_model = DecisionTreeClassifier(random_state=1015)
xgboost_model = XGBClassifier(random_state=1015, eval_metric="logloss")
random_forest_model = RandomForestClassifier(random_state=1015)

# Evaluate each model
evaluate_model(decision_tree_model, "Decision Tree", X_train, X_test, y_train, y_test)
evaluate_model(xgboost_model, "XGBoost", X_train, X_test, y_train, y_test)
evaluate_model(random_forest_model, "Random Forest", X_train, X_test, y_train, y_test)
```

## --- Decision Tree Results ---

## Training Metrics:

Accuracy: 1.0000

Precision: 1.0000

Recall: 1.0000

F1 Score: 1.0000

## Test Metrics:

Accuracy: 0.8853

Precision: 0.7183

Recall: 0.7675

F1 Score: 0.7421

## Training Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	17838
1	1.00	1.00	1.00	4961
accuracy			1.00	22799
macro avg	1.00	1.00	1.00	22799
weighted avg	1.00	1.00	1.00	22799

## Test Classification Report:


	precision	recall	f1-score	support
0	0.94	0.92	0.93	4474
1	0.72	0.77	0.74	1226
accuracy			0.89	5700
macro avg	0.83	0.84	0.83	5700
weighted avg	0.89	0.89	0.89	5700

## Training Confusion Matrix:

```
[[17838    0]
 [    0 4961]]
```

## Test Confusion Matrix:

```
[[4105  369]
 [ 285  941]]
```



### --- XGBoost Results ---

#### Training Metrics:

Accuracy: 0.9580  
 Precision: 0.9943  
 Recall: 0.8117  
 F1 Score: 0.8938

#### Test Metrics:

Accuracy: 0.9360  
 Precision: 0.9556  
 Recall: 0.7365  
 F1 Score: 0.8319

#### Training Classification Report:

	precision	recall	f1-score	support
0	0.95	1.00	0.97	17838
1	0.99	0.81	0.89	4961
accuracy			0.96	22799
macro avg	0.97	0.91	0.93	22799
weighted avg	0.96	0.96	0.96	22799

#### Test Classification Report:


	precision	recall	f1-score	support
0	0.93	0.99	0.96	4474
1	0.96	0.74	0.83	1226
accuracy			0.94	5700
macro avg	0.94	0.86	0.90	5700
weighted avg	0.94	0.94	0.93	5700

#### Training Confusion Matrix:

```
[[17815   23]
 [  934 4027]]
```

#### Test Confusion Matrix:

```
[[4432   42]
 [ 323  903]]
```



--- Random Forest Results ---

Training Metrics:

Accuracy: 1.0000

Precision: 1.0000

Recall: 0.9998

F1 Score: 0.9999

Test Metrics:

Accuracy: 0.9335

Precision: 0.9639

Recall: 0.7178

F1 Score: 0.8228

Training Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	17838
1	1.00	1.00	1.00	4961
accuracy			1.00	22799
macro avg	1.00	1.00	1.00	22799
weighted avg	1.00	1.00	1.00	22799

Test Classification Report:

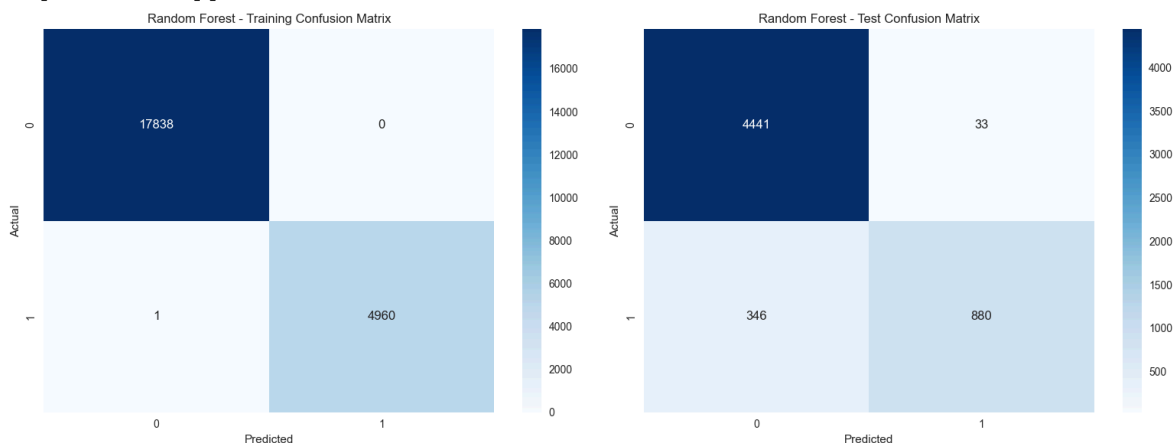
	precision	recall	f1-score	support
0	0.93	0.99	0.96	4474
1	0.96	0.72	0.82	1226
accuracy			0.93	5700
macro avg	0.95	0.86	0.89	5700
weighted avg	0.94	0.93	0.93	5700

Training Confusion Matrix:

```
[[17838    0]
 [    1 4960]]
```

Test Confusion Matrix:

```
[[4441   33]
 [ 346  880]]
```



## 2.3 Applying SHAP to XGBOOST



SHAP (SHapley Additive exPlanations) provides a principled way to interpret machine learning models by assigning each feature a contribution value for individual predictions, based on game theory.

- In this code, `summary_plot (bar)` shows which features have the greatest average impact across all predictions, while the default `summary_plot (dot)` reveals how feature values (high vs. low) influence predictions.
- The waterfall plot breaks down one prediction, showing how each feature nudges the output up or down from the model's baseline.
- Finally, the scatter plots (dependence plots) highlight how individual feature values affect their SHAP contributions, revealing trends and potential interactions.

```
In [84]: import shap

# SHAP Analysis
# Create a SHAP explainer object
explainer = shap.Explainer(xgboost_model)

# Calculate SHAP values
shap_values = explainer(X_test)

# Summary plot of SHAP values
shap.summary_plot(shap_values, X_test, plot_type="bar")
shap.summary_plot(shap_values, X_test)

# SHAP force plot for a specific prediction (e.g., first test instance)
shap.plots.waterfall(shap_values[0])

# SHAP dependence plots for top features
# Get feature names (adjust this if your X_test doesn't have column names)
feature_names = (
    X_test.columns
    if hasattr(X_test, "columns")
    else [f"feature_{i}" for i in range(X_test.shape[1])]
)

# Plot dependence plots for top 3 features based on mean absolute SHAP va
top_features_idx = np.argsort(-np.abs(shap_values.values).mean(0))[:3]
for idx in top_features_idx:
    shap.plots.scatter(shap_values[:, idx], color=shap_values)
```

